The Ultimate Guide to
Performance Engineering:
Bottleneck Patterns,
Diagnostics, and Solutions
Across Oracle, Java, Linux, AWS
Microservices, and Distributed
Systems

Symptom	& Likely Bottleneck	¿ Linux/Java/Oracle Clues	💋 Action Plan
High CPU Usage (80–100%)	- Tight loop- GC Overhead- Inefficient code- Oracle parsing overhead	- top, pidstat -u: Java PID at top- top -H: one thread near 100%- jstack: RUNNABLE thread(s) in same method- AWR: High "parse CPU"	- Capture multiple jstack dumps- Optimize logic, eliminate loops- Use bind vars to reduce Oracle parse load- Use jfr or perf to trace CPU usage
High Memory Usage (steady increase)	- Memory leak- Cache overflow- Long-lived object graphs	- free -m: swap increasing- jmap -histo / jmap -dump- GC logs: frequent FGCs- Oracle PGA/SGA nearing OS limits	- Dump heap, analyze with MAT- Tune caches- Fix long-lived references- Tune Oracle SGA/PGA to leave headroom
High CPU + Memory	- GC Thrashing- Memory leak + promotion failures	- jstat -gcutil: OldGen > 90%- jstack: GC threads RUNNABLE- GC logs: frequent FGCs, CMS failures	- Tune GC (G1GC, CMS, ZGC)- Adjust heap size- Reduce allocation rate- Fix object retention
High I/O Wait (%wa)	- Disk bottleneck- Oracle table scan- Log/trace overload- NFS delay	- vmstat: high wa, low id- iostat -x: %util > 90%, await > 20ms- AWR: db file sequential/scattered read high	- Use SSDs- Optimize queries with proper indexes- Reduce disk logging- Check NFS/storage config
High Thread Count	- Thread leak- Blocking operations spawning threads	- ps -eLf: nlwp very high- jstack: many similar threads- ulimit -u close to max	- Cap thread pools- Use executor reuse- Identify thread leaks- Close file handles/sockets properly
Many BLOCKED Threads	- Lock contention- Synchronized hotspots- Oracle TX enqueue locks	- jstack: many BLOCKED on same monitor- One thread holding monitor- Oracle: enq: TX - row lock contention	- Refactor code for fine-grained locking- Use ConcurrentHashMap etc Identify hot rows in DB and avoid them
Threads WAITING/TIMED_WAITIN G	- DB/API slowness- I/O bound calls-	- jstack: threads in socketRead- Oracle: slow	- Use async I/O- Tune DB/APIs- Add timeout + circuit breaker- Use connection pooling

	Socket read	queries in AWR- Logs:	
	timeout	external call slowness	
	timeout	external can slowness	
GC Threads Using CPU	- Over-sized	- jstack: GC threads	- Tweak GC generation sizes-
	Eden-	RUNNABLE- jstat: high Eden +	Tune
	Promotion	frequent FGC- GC logs:	CMSInitiatingOccupancyFractio
	failure-	CMS/G1 failures	n or G1 region size- Switch to
	Memory	·	ZGC/G1 if using old collector
	churn		,
OOM Errors	- Memory	- Logs: OutOfMemoryError-	- Analyze heap dump
	leak- Poor	heapdump.hprof- jstat -gcutil:	(MAT/VisualVM)- Fix leaks- Tune
	heap sizing-	FGCTime spikes	JVM heap (-Xmx)- Free non-
	Unclosed		heap (buffers, file I/O)
	buffers/cache		
	S		
Load Avg > CPU Cores	- CPU	- sar -q: runq > cores- top:	- Profile CPU vs I/O source-
	pressure- I/O	high load- iostat: high await	Tune thread pools- Add more
	delay- Thread		compute nodes (scale out)
	pool		
	exhaustion		
Oracle Wait Events (high	- Slow I/O-	- AWR/ASH: db file seq/scat	- Add missing indexes- Move
wait time)	Row lock-	read, log file sync, latch free,	logs to SSD- Reduce commit
wait time;	Latch/cache	buffer busy-	rate- Tune SGA pools (library,
	issues- Redo	v\$session/v\$event	shared)
	bottlenecks	vysession/ vyevent	shareu)
	bottleflecks		
Connection Pool	- DB not	- Logs:	- Increase pool size (carefully)-
Starvation	responding-	ConnectionTimeoutException	Tune SQL for speed- Fix JDBC
	JDBC leaks-	- jstack: waiting on	leaks (unclosed connections)-
	Pool too small	getConnection()- Oracle:	Use pool monitoring (HikariCP
		v\$session shows maxed	etc)
		connections	·
Slow External Calls (APIs)	- DNS	- jstack: waiting on	- Use persistent HTTP
	resolution	HTTPClient/Socket- Logs:	connections- Pre-resolve DNS-
	delay- Remote	endpoint latency spikes- dig,	Use local cache- Tune HTTP
	API latency-	ping, curl -w slow	client (timeouts, TLS reuse)
	TLS		
	handshake		
	overhead		
DR Slow Quaries	Missing	AMP: Top SOL by planted	- Add or rebuild indexes- Use
DB Slow Queries	- Missing	- AWR: Top SQL by elapsed-	
	indexes- High	v\$sql shows high	bind vars for plan stability-
	cardinality		Analyze stats (DBMS_STATS)
	1	I	l .

	joins- Bad stats	elapsed_time- EXPLAIN PLAN is suboptimal	
High GC Frequency High Redo/Log Write Wait	- Short-lived object spam- Small young gen- Memory leak - Too frequent commits- Slow disk for	- GC logs: frequent Pause Young- jstat -gc shows frequent minor GC- Memory not recovered - AWR: log file sync, log file parallel write- Redo write time high	- Increase young gen- Pool objects- Use escape analysis- friendly patterns - Batch commits- Move logs to fast disk- Increase LOG_BUFFER
	redo- Low redo buffer		
Excessive Hard Parses	- Literal SQL- Poor cursor reuse- High parse CPU	- AWR: "parse count (hard)" high- library cache lock waits- v\$sqlarea has many similar SQLs	- Use bind variables- Set cursor_sharing=force (carefully)- Increase shared_pool_size

Symptom Symptom	Likely Bottleneck	Kinux/Java/Oracle Clues	Action Plan
High Context	- Thread	- pidstat -w: high cswch/s	- Reduce locking (use ReentrantLock,
Switches	contention- Lock	(>10k/sec)- vmstat: high	LongAdder)- Replace blocking I/O with
	starvation-	cs- jstack: many threads	async- Profile code for hot spots
	Blocking I/O	BLOCKED or WAITING	
	overload		
High Steal Time	- Virtualization	- mpstat: %steal > 5%-	- Contact cloud/infra team- Migrate
(%st)	CPU overcommit-	iostat: normal but app	instance- Isolate workloads
	Noisy neighbor	slow	
	VMs		
High Run Queue (r	- CPU saturation-	- vmstat: r > #vCPUs- top:	- Optimize threads, locks- Scale out
in vmstat)	Lock contention-	CPU near 100%	horizontally
	Thread starvation		
High Minor Page	- Memory	- pidstat -r: minflt/s very	- Pool/reuse objects- Tune JVM GC for
Faults	fragmentation-	high- vmstat: high pgfault	allocation patterns
	Frequent object		
	creation		

High Major Page	- Memory under-	- pidstat -r: majflt/s > 0-	- Increase RAM- Reduce heap size-
Faults	provisioned- Swap	vmstat: si/so > 0	Avoid swap-intensive workloads
Taults	-	VIII3tat. 31/30 > 0	Avoid swap intensive workloads
	usage		
High ksoftirqd CPU	- Network packet	- top: ksoftirqd at top-	- Tune NIC interrupt coalescing-
	storm- Interrupt	mpstat -I S: softirq high-	Distribute IRQs with irqbalance
	coalescing issues	/proc/softirgs counters	
	-		
High Disk Queue	- Disk under-	- iostat -x: avgqu-sz > 1-	- Add SSDs- Tune I/O scheduler (noop,
Depth	provisioned- IOPS	High await (>10ms)	deadline)
	saturation		
Buffer Cache	- Hot blocks- Poor	- AWR: buffer busy waits	- Rebuild/index hot tables- Use hash
Contention	index design	high- v\$bh: same block#	partitioning or reverse keys
(Oracle)	illucx uesign	=	partitioning of reverse keys
(Oracle)		hotspots	
Redo Log	- Frequent small	- AWR: log file sync wait	- Batch commits- Move redo logs to
Contention	commits- Slow disk	event high	SSD/NVMe- Tune LOG_BUFFER
(Oracle)	for redo logs		
	- 1 1 1/2		
High log file	- Redo log I/O	- AWR: log file parallel	- Separate redo logs from datafiles- Fast
parallel write	bottleneck	write wait high	storage for redo logs
Oracle eng: TX -	- Index block hot	- AWR: TX - index	- Rebuild index- Use reverse key/index
index contention	spots	contention wait- Hot	partitioning
	•	index blocks	
Library Cache	- Hard parsing-	- AWR: library cache:	- Use bind variables- Increase
Mutex Contention	Frequent SQL	mutex X waits- v\$sqlarea:	shared_pool_size
	invalidation	many unique SQLs	
Checkpoint I/O	- Log switches too	- AWR: checkpoint	- Increase redo log size- Tune
Spikes (Oracle)	frequent-	complete messages- log	checkpoint intervals
Spines (Gracie)	Checkpoint tuning	file switch timeouts	checkpoint intervals
	issues	The switch timeouts	
	133003		
File Descriptor	- Resource leak	- lsof -p <pid>: open files</pid>	- Close file handles in code- Raise nofile
Exhaustion	(unclosed	near ulimit -n- ulimit -n	limit
	files/sockets)- High	low	
	concurrency		
TI 10 1	- " '	10.437 A 1:! 1	
Thread Pool	- Too small pool-	- JMX: ActiveThreads ==	- Increase pool size- Use async patterns-
Exhaustion (App	Blocking I/O-	MaxThreads- Thread	Tune DB/API calls
Layer)	Downstream	dump: many WAITING	
	slowness		
GC Time Spikes	- Tenuring	- GC logs: long Pause Full-	- Tune MaxTenuringThreshold- Pre-size
- c opineo	threshold too low-	jstat: OldGen spikes	young gen- Analyze object allocation
		J. 200 - 100	,

	Large object allocation		
Connection Leaks (DB/JDBC)	- Unclosed connections- No timeout settings	- HikariCP/DBCP metrics: active == max- Oracle: many inactive sessions	- Enforce try-with-resources- Set timeouts- Leak detection (leakDetectionThreshold)
Oracle latch: row cache objects	- Heavy DDL- Frequent object invalidation	- AWR: latch: row cache objects waits	- Reduce DDL at runtime- Use static object structures
Excessive log file switch (archiving needed)	- Archiving backlog- Small redo logs	- Alert log: frequent switches- v\$log: status ACTIVE/CURRENT	- Increase redo log size- Tune archive destinations
High Network Retransmissions	- Packet loss- NIC issues	- netstat -s: retransmits > 0- ifconfig: RX/TX errors	- Check cables, NICs- Adjust MTU- Optimize TCP stack
Frequent ORA- 01555: Snapshot Too Old	- Undo retention too low- Long- running queries	- AWR: ORA-01555 errors- v\$undostat	- Increase UNDO_RETENTION- Break large queries into chunks
High TCP TIME_WAIT Count	- Short-lived connections- No keep-alive reuse	- ss -s: TIME_WAIT > 10k- netstat -an	- Enable keep-alive- Use connection pooling- Tune tcp_fin_timeout
High Oracle cursor: pin S Waits	- Frequent hard parses- SQL not shared	- AWR: cursor: pin S wait on X	- Use bind vars- Tune cursor sharing- Increase shared_pool_size
Network Latency Spikes	- External API slowness- DNS resolution delays	- ping, mtr, traceroute to APIs- App logs: response time high	- Use persistent connections- Pre- resolve DNS- CDN/proxy caching
High Java Finalizer Activity	- finalize() abuse- Unreleased resources	- jstat: high FGC- Thread dump: Finalizer threads RUNNABLE	- Remove unnecessary finalizers- Use try-with-resources- Consider PhantomReferences
Long GC Stop-the- World Pauses	- Large heap + Full GC- CMS fragmentation	- GC logs: "Pause Full" > 500ms- jstat: high FGC	- Switch to G1GC/ZGC- Reduce OldGen occupancy- Tune CMSInitiatingOccupancyFraction
High SoftIRQ CPU Usage	- Network flood- Interrupt handling bottleneck	- mpstat -I S: high softirq- /proc/softirqs counters	- Use irqbalance- Offload NIC features (GRO/TSO)

Oracle AWR Analysis – Bottleneck Patterns & Resolutions

Symptom/Pattern	S Likely	⊘ How to Detect	Resolution Plan
(AWR Clues)	Bottleneck	(AWR/ASH)	
Top SQL by Elapsed Time / Gets	- Expensive queries- Missing indexes- Large table scans	- AWR SQL ordered by elapsed_time / buffer_gets-High disk_reads, rows_processed	- Optimize SQL (joins, predicates)- Add/rebuild indexes- Use partitioning
High db file sequential read	- Slow single-block I/O- Index scans on large tables	- AWR: db file sequential read top wait- Top SQL: high disk_reads	- Check IOPS/disk latency- Add indexes- Tune storage layer
High db file scattered read	- Full table scans- Poor SQL design	- AWR: db file scattered read wait- Top SQL: full scans (rows_processed high)	- Rewrite SQL- Add indexes- Avoid full scans where possible
High buffer busy waits	- Hot blocks- Block contention	- AWR: buffer busy waits high- Segment statistics: hot objects	- Reorganize data- Use hash partitions/reverse key indexes- Tune block sizes
High log file sync	- Frequent commits- Redo log I/O bottleneck	- AWR: log file sync in top waits- Average sync time > 5ms	- Batch commits- Optimize redo logs (size, placement)- Tune storage IOPS
High log file parallel write	- Slow redo log writes- Storage I/O limits	- AWR: high log file parallel write wait- High redo write times	- Use fast disks (SSD/NVMe)- Tune log buffer size- Reduce commit frequency
High free buffer waits	- DBWR unable to keep up- Slow disk writes	- AWR: free buffer waits in top waits- High dirty buffer count	- Increase DBWR I/O capacity- Add disks / use faster storage- Check db_writer_processes
High library cache lock/mutex	- Cursor contention- Frequent hard parses	- AWR: library cache waits- High parse count in AWR	- Use bind variables- Increase shared pool- Reduce hard parsing SQLs
High latch: cache buffers chains	- Hot block in buffer cache	- AWR: latch: cache buffers chains- Segment stats: hot object	- Reorganize data- Partition tables- Tune SQL to avoid hot blocks
High row cache lock	- Data dictionary contention- DDL during runtime	- AWR: row cache lock in top waits	- Avoid DDL during peak- Tune shared_pool_size
High enq: TX - row lock contention	- Hot rows- Update contention	- AWR: enq: TX in top waits- Blockers in v\$session	- Optimize transactions- Reduce row-level contention- Use

			SELECT FOR UPDATE SKIP
			LOCKED
High enq: TX - index	- Hot index blocks	- AWR: enq: TX - index waits	- Rebuild index- Use reverse
contention			key/index partitioning
High db file parallel	- Slow datafile	- AWR: db file parallel write	- Optimize disk layout- Tune
write	writes- High I/O	top event- I/O timing >	dbwr processes- Use faster disks
	workload	10ms	
direct path	- Direct I/O	- AWR: direct path read or	- Optimize bulk operations- Use
read/write high	operations (e.g.,	write top waits	parallel DML carefully
	large queries, bulk		
	loads)		
latch free high	- Latch contention	- AWR: latch free top wait-	- Tune latch-intensive code-
	(shared resource	v\$latch stats high	Reduce shared resource usage
	hot spot)		
High Parse CPU Usage	- Literal SQL-	- AWR: high parse time CPU-	- Use bind variables- Tune cursor
	Excessive parsing	High parse count (hard)	sharing (FORCE, SIMILAR)
High cursor: pin S	- Pin contention in	- AWR: cursor: pin S wait-	- Use bind variables- Tune shared
Waits	library cache	High mutex waits	pool size
High db file	- Inefficient index	- AWR SQL plan shows many	- Rebuild index- Consider index
sequential read in	access- B-Tree	index leaf blocks	compression or partitioning
Index Reads	depth too high		
High undo segment	- Long transactions-	- AWR: undo segment waits-	- Increase UNDO_RETENTION-
waits	Undo retention too	v\$undostat metrics	Commit frequently
	low		
High gc Waits in RAC	- Global cache	- AWR: gc cr request, gc	- Optimize interconnect
	transfer latency	buffer busy waits- High	(InfiniBand, RoCE)- Tune cache
		interconnect latency	fusion settings
SQL Execution Time >	- Suboptimal	- AWR: SQL with high	- Gather fresh stats- Use SQL
Expected	execution plans-	elapsed time- v\$sql_plan	Plan Baselines- Tune optimizer
70	Stale stats	shows bad plan	parameters
cell single block	- Smart scan not	- AWR: cell single block top	- Enable smart scans (direct path
physical read	used	event- Exadata-specific	reads)- Use full scans when
(Exadata)		metrics	beneficial
<u>L</u>	1		

Cache-Related Performance Bottlenecks – Advanced Guide

Cache Layers Involved

<u>Layer</u>	Cache Type	Example
Application Layer	In-memory caches	Ehcache, Guava, Caffeine, Hazelcast
JVM Layer	Code cache, object heap	JIT-compiled methods, permgen/metaspace
Oracle DB	Buffer cache, shared pool, library cache, dictionary cache	Data blocks, SQL cursors, PL/SQL, data dictionary
OS Layer	Filesystem cache, page cache	Linux page cache, disk read-ahead
Network Layer	DNS cache, connection pools	DNS resolver cache, HTTP keep-alives

Q Cache Bottleneck Patterns & Resolutions

Symptom	Likely Cache Issue	Detection/Diagnostics	<page-header> Resolution Plan</page-header>
High GC/Memory	- Cache too large	- jmap -histo shows large	- Configure max size/TTL
Usage in Java App	(unbounded growth)	cache classes	- Use
	- Inefficient eviction	- jstat -gc shows OldGen	LinkedHashMap/Caffeine with
		growth	eviction
		- JVM metrics: cache size	- Tune eviction policies (LRU,
			LFU)
5 15 11 66		661 (15.11.66	
Frequent Full GCs	- Cache churn creating	- GC logs: frequent Full GC	- Reduce cache churn
	objects	- jstat -gcutil: high OldGen	- Pool objects
	- Poor object pooling		- Pre-warm caches
Class Overies in	- Low buffer cache hit	- AWR: Buffer Cache Hit	Ingressed by each a size
Slow Queries in			- Increase db_cache_size
Oracle	ratio	Ratio < 90%	- Tune SQL for index usage
	- SQLs causing physical	v\$sysstat/v\$sysmetric:	- Keep hot tables in memory
	1/0	cache misses	
Library Cache Latch	- Excessive hard parsing	- AWR: library cache	- Use bind variables
Contention (Oracle)	- Literal SQL (no bind	lock/library cache mutex	- Increase shared_pool_size
Contention (Oracle)		•	_ = =
	vars)	waits	- Avoid dynamic SQL in loops
		- v\$sqlarea: many similar	
		SQLs	

Shared Pool	- SQLs > 4KB	- AWR: high parse time CPU	- Increase shared_pool_size
Fragmentation (Oracle)	- Poor cursor sharing	- v\$sgastat: free memory low in shared pool	- Use CURSOR_SHARING wisely - Pin PL/SQL packages
High buffer busy waits (Oracle)	- Hot blocks in buffer cache - Contention on frequently accessed data	- AWR: high buffer busy waits - Segment statistics: hot objects	Reorganize tablesUse hash partitioning or reverse key indexes
Poor File System Cache Performance	- Inefficient OS page cache usage - Double buffering	- vmstat: high page-in/out - sar -B: high page faults - iostat: high disk reads	- Use direct_io (DB side)where applicable- Tune Linux dirty ratios(/proc/sys/vm/dirty_*)
DNS Resolution Delays	- DNS cache miss at app layer - Frequent external lookups	dig, nslookup: slowresponseLogs: connection delays onfirst API call	Use persistent connectionsPre-resolve DNS entriesLocal DNS caching (e.g., dnsmasq)
High Connection Pool Churn	- No connection caching - Low pool sizes	- Connection pool metrics: high borrowed but no reuse - Logs: frequent connection opening/closing	 Use proper pooling (HikariCP, c3p0) Set minIdle, maxPoolSize Reuse connections across requests
Cache Invalidation Storm	- Sudden eviction of large portions of cache (e.g., cache.clear())	- Logs: bulk cache evictions - Spikes in DB or backend calls after	Use selective invalidation (by key)Staggered TTLs ("jittered cache")Versioned cache keys
Cache Stampede Effect	- Many threads request same data simultaneously (cold cache)	- Logs: many requests for same key after expiry - Spikes in backend load	- Add request coalescing("single-flight" logic)- Use per-key locking- Stagger cache expirations
Hot Key Contention	- One key heavily accessed by all threads (e.g., a global config)	 Cache metrics: hot key hit rate disproportionate Thread dump: threads waiting on cache fetch 	- Shard data if possible - Add randomization to keys - Use per-tenant cache
Cache Eviction Too Aggressive	- Low max size / short TTL - Premature eviction of useful data	- Metrics: low hit rate despite high cache usage - Logs: frequent refetch from DB	Increase cache sizeAdjust TTLsReview eviction policy (LRU vs. LFU)

Oracle latch: cache	- Hot block access in	- AWR: latch: cache buffers	- Partition or hash table
buffers chains	buffer cache	chains top wait	- Reduce block-level
Waits		- Segment stats: hot object	contention
SGA Memory	- SGA too small for	- AWR: high free memory	- Increase SGA (sga_target)
Pressure (Oracle)	workload	fragmentation	- Use sga_target vs. manual
	- Fragmented pools	- v\$sgastat: low free memory	sizing based on workload
High cache buffer	- Large block sizes	- AWR: latch misses > 0	- Use smaller block size if
chains latch misses	- Hot objects	- Hot segment analysis	possible
			- Optimize schema design

Advanced Debugging Techniques for Cache Issues

Tool/Command	Purpose Purpose	Example
jmap -histo	Find largest objects in	jmap -histo:live <pid></pid>
	JVM	
jstack	Find threads blocked on	Look for BLOCKED on cache class
	cache	
jstat -gcutil	JVM GC trends (churn)	jstat -gcutil <pid></pid>
AWR SQL ordered by Gets	High cache usage SQLs	Check top SQL by buffer_gets
v\$sysstat	Buffer cache hit ratios	select name, value from v\$sysstat where name like
		'%cache%';
sar -B, vmstat	OS cache activity	sar -B for page faults, vmstat for si/so
lsof, ulimit -n	Check file descriptor leaks	lsof -p <pid></pid>
dtrace/perf	Trace kernel-level caching	Use for syscall-level I/O
Cache Metrics (Caffeine,	Monitor hit/miss rates	Expose via JMX or logs
Guava)		

AWS Cloud Microservices Performance Issues – Deep Dive

	Likely Bottleneck	@ Detection Methods	Resolution Strategy
High Latency in API Gateway/ALB	- Backend slow responses - Cold starts (Lambda) - TLS handshake delays	- AWS X-Ray traces: high API latency - ALB target response times - CloudWatch metrics: TargetResponseTime, 5XX errors	- Enable caching in API Gateway - Use provisioned concurrency for Lambda - Optimize TLS handshake (TLS 1.3, keep-alives) - Reduce backend response times
Container CPU Spikes (ECS/EKS)	- Inefficient app code - Thread leaks - GC pauses	- CloudWatch: CPUUtilization spikes - kubectl top pods in EKS - JVM jstat + GC logs	- Optimize code/hotspots - Adjust CPU limits/requests - Tune GC (G1GC, ZGC)
Container Memory Usage Spikes	- Memory leaks - Unbounded cache growth - Large object creation	- CloudWatch: MemoryUtilization trends- JVM heap usage (Prometheus, JMX)- Container OOM kills	- Set memory limits + requests - Profile heap usage (jmap, VisualVM) - Fix leaks, cache sizing
Cold Start Latency (Lambda)	- Unprovisioned Lambda - Large package size - VPC-enabled Lambda	- CloudWatch: Duration, InitDuration metrics - X-Ray traces	- Use provisioned concurrency - Minimize deployment package size - Use Lambda SnapStart (Java) - Avoid VPC unless necessary (or use ENI reuse)
High API 5XX Errors	- Backend failures - Rate limits - API Gateway integration timeouts	- CloudWatch: 5XXError metrics - X-Ray: backend service faults	- Set retries with backoff - Handle errors gracefully - Optimize downstream services

High Latonay Variability	Noisy poichbor	CloudWatch: DOOL atomor anilyan	Uso multi 47
High Latency Variability	- Noisy neighbor	- CloudWatch: P99Latency spikes	- Use multi-AZ
(P99 Spikes)	effect	- VPC Flow Logs: high retransmits	deployments
	- Network	- JVM GC logs	- Tune GC for low-
	bottlenecks		latency
	- GC pauses		- Use HTTP/2
			multiplexing
Throttling in AWS	- Exceeding AWS	- CloudWatch: Throttles metrics	- Implement client-
Services (429 Errors)	API limits	- AWS SDK retries/exceptions	side throttling
	- Unthrottled burst		- Use service quotas
	traffic		APIs to monitor
			- Stagger workloads
High DB Latency	- Query inefficiency	- RDS Performance Insights: Top SQLs	- Optimize SQL
= -			(Aurora): indexes,
(Aurora/DynamoDB)	- Hot partitions	- DynamoDB CloudWatch:	,
	- IOPS saturation	ThrottledRequests	partitioning
			- DynamoDB: Use
			partition keys,
			adaptive capacity
			- Scale read replicas
			(Aurora)
Slow Microservice-to-	- Network latency	- X-Ray: high service latency	- Enable HTTP/2/gRPC
Microservice Calls	- TLS handshake	- dig, mtr, curl -w tests	- DNS caching
	overhead	ang, many carrier to costs	- Enable TLS session
	- DNS resolution		reuse
	delays		reuse
	uelays		
Service Discovery Delays	- DNS TTL too low	- kubectl logs coredns in EKS	- Increase DNS TTLs
	- Overloaded	- High DNS resolution time in X-Ray	- Scale CoreDNS
	CoreDNS		replicas
			- Use service mesh
			sidecar DNS caching
High Inter-Zone Latency	- Cross-AZ traffic	- VPC Flow Logs: cross-AZ traffic	Hea cama A7
night inter-zone Latency		_	- Use same AZ
	- Unoptimized	- X-Ray traces: high latency spans	placement where
	routing		possible
			- Optimize inter-
			service
			communication
Frequent Pod Evictions	- Node resource	- CloudWatch: Evictions metric	- Set proper
(EKS)	exhaustion	- kubectl describe pod	CPU/memory
	- Misconfigured	·	requests
	resource limits		- Use Cluster
			Autoscaler
			- Right-size nodes
			MgHC-312E HOUES

Lambda API Gateway	- Lambda	- CloudWatch: IntegrationLatency	- Optimize Lambda
Timeouts	exceeding API	near 29s	· ·
Timeouts	_		logic
	Gateway 29s limit	- X-Ray traces	- Break large tasks
			into smaller functions
			- Use Step Functions
			for orchestration
Cache Miss Spikes	- TTL expiries	- CloudWatch: CacheHits/Misses	- Increase TTLs
(ElastiCache/DAX)	- Hot key eviction	metrics	- Use distributed
		- Application logs	locking (e.g., Redlock)
			- Avoid sudden bulk
			invalidations
11. 1 EDG 1 .	D. LIODS I		
High EBS Latency	- Disk IOPS limits	- CloudWatch: VolumeQueueLength,	- Use provisioned
(GP2/GP3)	- Burst balance	VolumeThroughputPercentage	IOPS (io1/io2)
	depletion		- Monitor burst
			balance
			- Distribute workload
			across volumes
High Lambda	- Traffic spike	- CloudWatch: ConcurrentExecutions	- Set reserved
Concurrent Executions	- No reserved	metric	concurrency
	concurrency	- Throttling errors in logs	- Use weighted
			routing for canary
			deployments
			- Enable async event
			sources
VDC Notwork	High FMI usage	VDC Flow Loggy retropemits, high	Lico Nitro basad
VPC Network	- High ENI usage	- VPC Flow Logs: retransmits, high	- Use Nitro-based
Bottlenecks	- Insufficient	latency	instances
	bandwidth	- CloudWatch: Network metrics	- Scale horizontally
			- Enable Enhanced
			Networking (ENA)
Service Mesh Overhead	- Sidecar	- CloudWatch: sidecar metrics	- Optimize Envoy
(App Mesh/Istio)	CPU/memory	- Istio metrics: pilot, envoy stats	configs
	usage		- Reduce mTLS if not
	- mTLS encryption		needed for internal
	cost		traffic
			- Offload metrics
			collection
X-Ray Trace Data	- Sampling too low	- No traces for requests in X-Ray	- Adjust sampling
Missing	- Misconfigured	and the state of t	rules
1411331115	SDKs		- Instrument SDKs
	201/2		
			properly

			- Use OpenTelemetry
			for custom metrics
High DynamoDB	- Hot partitions	- CloudWatch: ThrottledRequests,	- Use adaptive
Throttling	- Exceeded	Provisioned Read Capacity Units	capacity
	RCU/WCU		- Distribute keys
			- Switch to on-
			demand mode
Inefficient Autoscaling	- Cooldown delays	- CloudWatch: delayed scale-in/scale-	- Tune scaling policies
Response	- Metric lag	out	- Use predictive
			scaling
			- Leverage step
			scaling with proper
			thresholds

Advanced Tools for AWS Microservices Performance Analysis

√ Tool	Purpose	Example
AWS X-Ray	Distributed tracing	Find service-level bottlenecks
CloudWatch Logs Insights	Query logs at scale	Analyze latency spikes
CloudWatch Metrics	Infrastructure + app metrics	CPU, memory, latency, errors
AWS Distro for OpenTelemetry	Custom metrics + tracing	Ingest into CloudWatch/Prometheus
VPC Flow Logs	Network diagnostics	Detect packet loss, latency
kubectl top, kubectl logs	EKS performance	Resource usage per pod
JMX Exporter + Prometheus	JVM metrics	GC, heap, thread pool analysis
Load Testing (k6, JMeter)	Synthetic load	Test under production-like conditions
CloudTrail	API activity tracking	Detect throttles, access patterns

Additional AWS Cloud Microservices Performance Bottlenecks

Hidden/Advanced Patterns

S Bottleneck	S Details	Detection & Tools	Resolution Strategy
API Gateway	- Large HTTP headers	- ALB logs: ClientClosedConnection errors	- Minimize
Header Bloat	cause latency	- CloudWatch 5XX spikes	JWT/claims size
	- Custom headers		- Compress headers

Lambda Execution Environment Leaks	(JWTs, tracing IDs) can cross 8KB+ - Static variables persist across invocations, leading to state leaks	- Unpredictable behavior in cold vs warm starts - Hard-to-debug issues in stateful logic	where possible - Move metadata to payload - Avoid static mutable state - Use idempotent functions - Isolate logic in handlers
Excessive	- High log volume can	- Delayed CloudWatch metrics/alarms	- Filter logs (reduce
CloudWatch Logs	throttle ingestion	- High CloudWatch ThrottledLogEvents	verbosity)
Latency	- Delayed metric		- Use Kinesis
	generation		Firehose for high-
			throughput logs
Unbalanced Load in	- Sticky sessions or	- TargetResponseTime high for subset of	- Use round-robin or
ALB/EKS	DNS caching causes	pods	least-connections
	traffic skew	- X-Ray: some pods overloaded	- Use service mesh
			load balancing
Intermittent	- AWS backbone	- VPC Flow Logs: retransmits, high RTT	- Use Global
Network Latency	congestion	- Ping/traceroute: latency spikes	Accelerator for low-
(AWS Global	- Cross-region		latency routing
Network)	replication delays		- Localize traffic
			where possible
Lambda Package	- Large dependencies	- Deployment errors, slow cold starts	- Use Lambda layers
Size Bloat	inflate package	- S3 deployment requirement	- Minimize
	>250MB		dependencies
	XII		- Tree-shake unused
			libraries
EBS Burst Balance	- GP2 volumes	- CloudWatch: BurstBalance metric = 0	- Switch to
Exhaustion	exceed burst IOPS	- High VolumeQueueLength	GP3/io1/io2 with
			provisioned IOPS
			- Distribute load
			across volumes
Undersized NAT	- NAT GW bandwidth	- Slow outgoing traffic	- Use multiple NAT
Gateway	bottleneck (~5 Gbps	- CloudWatch NAT metrics: BytesOut at	GWs per AZ
	limit)	limit	- Use Egress-Only
			Internet Gateway if
			IPv6

Unantimized Star	Long rupping states	Ston Functions matrice: Fysicition Time	Split long states
Unoptimized Step Functions Execution	- Long-running states cause timeouts - High state transitions cost	- Step Functions metrics: ExecutionTime, BilledDuration	- Split long states - Reduce state transitions - Consider Lambda batching
Excessive SNS/SQS Fan-Out Latency	- Many subscribers cause delivery delays	- CloudWatch: NumberOfMessagesSent high, delayed NumberOfMessagesReceived	- Use FIFO queues to preserve order - Consider Kinesis if high fan-out required
Noisy Neighbor in Multi-Tenant K8s Cluster	- One namespace hogs resources	- EKS metrics: node CPU/mem skewed - kubectl top pods imbalance	- Enforce resource quotas - Use LimitRanges per namespace
Container Image Bloat	- Large images (>1GB) cause slow pulls	- EKS: ImagePullBackOff errors - Long pod startup times	- Use minimal base images (distroless, alpine) - Multi-stage Docker builds
High Service Mesh CPU/Memory (Envoy)	- Sidecar proxies consume significant resources	- Prometheus:envoy_cluster_upstream_rq_time spikes- CPU/memory usage of sidecars	- Optimize Envoy config - Consider eBPF- based alternatives (Cilium)
Inefficient VPC Peering Architecture	- Cross-VPC traffic via peering scales poorly	- VPC Flow Logs: high cross-VPC traffic - Latency variance in cross-service calls	- Consider Transit Gateway or PrivateLink
RDS Connection Pool Saturation	- High connection churn from Lambda/ECS	- RDS metrics: DBConnections at max - App logs: Too many connections errors	- Use RDS Proxy - Reuse connections - Tune pool size per workload
Inefficient Data Fetch in Microservices	- N+1 query problem across services	- X-Ray: repetitive service calls - Logs: many API calls for single request	- Batch APIs - GraphQL federation - CQRS patterns
Overly Chatty Microservices	- High inter-service call rates cause latency	- X-Ray: many small calls - Network metrics: high packet rates	- Aggregate API calls - Use event-driven design (SNS/SQS, EventBridge)

Elastic IP Address	- Failover IP takes	- Failover tests: EIP takes ~60s to reattach	- Use Route 53
Reassign Delays	time to propagate		health checks +
			multi-value answers
			- Use Global
			Accelerator for low-
			latency failover
CloudWatch Alarms	- Delayed metrics	- Alarm triggered after >5 min delay	- Use high-resolution
Lag	cause late alerts		metrics (1s
			granularity)
			- Leverage real-time
			monitoring tools
			(Datadog, New
			Relic)

Thread Dump Analysis – The Ultimate Guide

Thread Dump Basics

- A Thread Dump is a snapshot of all JVM threads, showing:
 - Thread state: RUNNABLE, WAITING, BLOCKED, TIMED_WAITING
 - Stack traces: What the thread is doing
 - o Locks held/waiting on: monitor, ownable synchronizer
 - o **Thread names**: http-nio-8080-exec-1, ForkJoinPool.commonPool-worker-3
 - Daemon/Non-daemon status

When to capture?

- During high CPU usage
- During latency spikes
- When blocked threads suspected
- During memory leaks (GC threads active)

✓ How to capture?

- jstack <pid> (Linux/Unix)
- kill -3 <pid> → Outputs to stdout
- JVM tools: VisualVM, Java Mission Control, JFR
- Cloud: ECS/EKS sidecar jstack, AWS Lambda SnapStart diagnostics

Thread States & What They Indicate

Thread State	Meaning	Possible Bottleneck
RUNNABLE	Actively executing on CPU	Hot code path, tight loops, infinite loops, GC overhead
		de overneau
WAITING /	Waiting for external events (I/O,	Network/API delays, DB calls, thread pool
TIMED_WAITING	locks, sleeps)	starvation
BLOCKED	Waiting for a monitor/lock	Lock contention, deadlocks, synchronized
		bottlenecks
TERMINATED	Completed thread	Normal or unexpected (abrupt termination)
GC Threads	JVM GC activity	High CPU, GC pressure, object churn

\mathbb{Q} Bottleneck Patterns from Thread Dumps

Pattern	What to Look For	Action Plan
CPU Bottleneck (High RUNNABLE Threads)	Many threads in RUNNABLE at top of dump, in same code (e.g., parsing, loops, serialization)	Optimize code, algorithm, data structures; Profile with - XX:+UnlockDiagnosticVMOptions -XX:+PrintCompilation; Use perf, jfr
Lock Contention (BLOCKED Threads)	Many threads BLOCKED on same lock/monitor; One thread holding the lock	Refactor locking code; Reduce critical sections; Use ReentrantLock, StampedLock; Consider lock-free structures (ConcurrentHashMap, LongAdder)
Thread Pool Starvation	All threads in pool in WAITING; No available worker threads	Increase thread pool size; Analyze task execution time; Consider queue capacity
External I/O Blocking	Threads WAITING in socketRead, read0, poll, accept	Add timeouts; Use async/non- blocking I/O (NIO, Netty); Cache external calls
Deadlock Detected	Found one Java-level deadlock in dump	JVM reports deadlock with threads and locks; Break cycle; Redesign lock acquisition order

GC Overhead	GC threads in RUNNABLE; frequent GC logs	Tune heap size, GC algorithm;
		Reduce object allocation rate
ForkJoinPool/Parallel	ForkJoinPool worker threads in BLOCKED/WAITING	Reduce parallelism; Avoid nested
Streams Saturation		parallel streams; Control task size
High Thread Count	Thousands of threads; common prefix (e.g., http-	Check thread creation in code;
(Leak)	nio-8080 or Executor-)	Limit thread pools; Use close() on
		resources
Thread Waiting on	Threads in WAITING in JDBC driver calls (e.g.,	Optimize SQL; Use connection
Database Calls	socketRead)	pool; Check DB performance
		(AWR, slow queries)
Async Task Backlogs	Threads in WAITING in	Monitor queues; Increase pool
	ScheduledThreadPoolExecutor/CompletableFuture	sizes; Apply backpressure
		strategies
Finalizer/Reference	Threads in ReferenceHandler, Finalizer	Avoid finalize(); Use
Handler Threads		AutoCloseable; Reduce
Active		phantom/weak references

Typical Stack Frames in Thread Dumps

Method	What It Indicates	Related Bottleneck
java.net.SocketInputStream.socketRead0	Waiting for I/O	Network API delays
sun.nio.ch.EPollArrayWrapper.epollWait	I/O multiplexing wait	Normal; may indicate thread pool starvation if too many
java.util.concurrent.ThreadPoolExecutor.getTask	Waiting for task	Normal unless pool too small
java.util.concurrent.locks.ReentrantLock\$NonfairSync	Locking bottleneck	Lock contention
java.lang.Thread.sleep	Sleeping (deliberate delay)	May cause latency spikes if overused
java.lang.Object.wait	Synchronized wait	Inter-thread communication or resource lock
org.apache.catalina.connector.Request.doRead	Servlet input blocking	HTTP request I/O bottleneck
org.hibernate.loader.Loader	Hibernate ORM data loading	N+1 query issues, lazy loading

% Advanced Tools for Thread Dump Analysis

Tool	Purpose	Usage
Thread Dump Analyzer (TDA)	Visual thread dump analysis	Highlight threads by state, identify locks
fastThread.io	Online thread dump viewer	Detect deadlocks, contention hotspots
jClarity Censum / GCEasy	GC + thread dump correlation	Visualize GC impact on threads
JFR (Java Flight Recorder)	Profiling + thread activity	Correlate CPU, threads, GC events
async-profiler + Flamegraphs	CPU hotspots	Flamegraph visualization of code execution

Thread Dump Analysis Workflow

- 1 Capture 3-5 dumps at 10s intervals during issue.
- 2 Filter by thread state: focus on RUNNABLE, BLOCKED.
- 3 Group by thread name prefix: e.g., http-nio, ForkJoin.
- 4 Identify hot code paths: same method in multiple threads.
- 5 Look for locks held/waited: detect contention or deadlocks.
- 6 Cross-check with GC logs, CPU usage, external systems.
- 7 Form hypotheses and test fixes in dev/stage.

GC Log Analysis – The Ultimate Guide for Java Performance

Why Analyze GC Logs?

GC logs reveal:

- ✓ Allocation pressure
- **✓** Memory leaks
- ✓ Pause times affecting latency
- ✓ Throughput vs. responsiveness trade-offs
- GC algorithm effectiveness
- They bridge the gap between code behavior and JVM internals:
 - GC logs = "heartbeat" of the JVM.

III GC Log Essentials: Key Fields to Watch

Field	What It Shows	Impact on Performance
Pause Time (Pause Young, Pause Full)	How long app is paused	Directly affects P99 latency

GC Type (Young, Old, Mixed)	Which generation collected	Helps identify which space causes issues
Before/After Heap	Heap size before/after GC	Reveals allocation rate & retention
Frequency of GCs	How often GC runs	Too frequent = high churn, too rare = risk of OOM
Promotion Rate	Rate of objects moving to OldGen	High rate → risk of OldGen fill-up
Allocation Rate	How much memory app allocates per second	High rates = GC pressure
Concurrent Phase Durations	Time spent in background GC tasks	Long phases = background work stalling app
Finalizer Activity	Finalizer threads usage	Excessive use → leaks, GC delays

© Common GC Bottleneck Patterns

℘ GC Pattern	Symptoms in Logs	Ø Root Cause & Fix
GC Thrashing	Many Pause Young eventsLow pause times but	High object churn in EdenObject
(Frequent Minor GCs)	very frequentOldGen stable	allocation too fast
Promotion Failure	Frequent Full GC (promotion failed)	OldGen too smallObjects live too
		long
Long GC Pauses (Full	Few but long Pause Full eventsApp freezes	Heap too largeCMS
GC)		fragmentationOldGen leaks
Concurrent Mode	concurrent mode failure in logs	CMS can't finish before OldGen fills
Failure (CMS)		
GC Overhead Limit	JVM exits with GC overhead limit exceeded	JVM spends >98% time in GC
Exceeded		
High GC CPU Usage	GC threads in RUNNABLE in thread dumpGC CPU	Too frequent GCsLarge heapsCMS
	>30%	scanning cost
Metaspace OOM	java.lang.OutOfMemoryError: Metaspace	Excessive class loadingClassloader
		leaks
Frequent G1 Mixed	Pause Mixed (G1 Evacuation Pause) events	OldGen has high live dataMixed
GCs	frequent	collections trigger often
Unusual Pause	Long tail latencies (high P99)	Some pauses 5-10x longer than
Distribution		median

GC Log Formats & Enabling Logging

JVM	GC Log Option	Notes
Java 8	-XX:+PrintGCDetails -Xloggc: <file></file>	Standard format
Java 9+	-Xlog:gc*:file= <file></file>	Unified Logging
Java 11+	-Xlog:gc*:stdout:time,uptime,level,tags	More granular
Enable GC Log Rotation	-XX:+UseGCLogFileRotation -XX:NumberOfGCLogFiles=5 - XX:GCLogFileSize=20M	Avoid huge logs

6 GC Algorithm-Specific Insights

G1GC Patterns

Q Pattern	Clue in Logs	Action
Excessive Mixed GCs	Pause Mixed every few seconds	Reduce -XX:InitiatingHeapOccupancyPercent
Pause Time Spikes	Long Pause Young/Pause Mixed	Tune -XX:MaxGCPauseMillis
High Allocation Rate	Eden fills up fast	Increase YoungGen size (via -XX:NewRatio or G1NewSizePercent)

ZGC/Shenandoah

Q Pattern	Clue	Action
Large Heap + Low Pause	Pauses <10ms at large heaps	Expected behavior—good!
Pauses >10ms	Unexpected long pauses	Check allocation rate, GC threads

GC Log Analysis Tools

Tool	What It Does	Usage
GCViewer	Visual GC graphing	Parse -Xloggc files
GCEasy.io	Online GC log analyzer	Upload logs, get insights
JClarity Censum	Commercial GC analyzer	Advanced analysis, not free

[Eclipse Memory Analyzer (MAT)]	Heap dump analysis	Find memory leaks
[JFR + JMC (Java Mission Control)]	Real-time GC + thread correlation	Enable with -XX:+FlightRecorder

- **Section** GC Tuning Workflow (Battle-Tested)
- **Capture GC Logs** → 10-30 mins under load
- Correlate with App Latency (CloudWatch, X-Ray, Prometheus)
- **✓** Identify Patterns:
 - GC frequency
 - Pause time distribution
 - OldGen/YoungGen occupancy
 - ✓ **Heap Dump at High GC** → Analyze leaks
 - **✓** Adjust GC Settings:
 - **G1**: MaxGCPauseMillis, InitiatingHeapOccupancyPercent
 - ZGC/Shenandoah: Heap size, thread count
 - **▼** Test Under Load → Rinse & repeat

UNIX Server Logs & Trace Analysis – Advanced SRE Guide

Key System Logs to Monitor

Log File	Path	Purpose Purpose Purpose
syslog / messages	/var/log/syslog or /var/log/messages	Kernel + general system messages
dmesg	dmesg or /var/log/dmesg	Boot time + hardware/kernel errors
auth.log	/var/log/auth.log	SSH, sudo, login failures
secure	/var/log/secure	PAM, security events
kern.log	/var/log/kern.log	Kernel errors, OOM killer, CPU stalls
journalctl	journalctl -xe	Combined systemd logs
application logs	/var/log/ <app>/</app>	App-specific or via journald
audit.log	/var/log/audit/audit.log	SELinux, permission audits
gc.log	Custom JVM GC log file	GC pause and allocation patterns
access.log, error.log	/var/log/nginx/, /var/log/httpd/	Web server latency, errors

System Performance Tracing: Core Commands

Command	Use Case	<u>Sample</u>	
top, htop	CPU/mem usage per process	top -H -p <pid> shows thread CPU</pid>	
vmstat 1	CPU, run queue, memory	Detect CPU saturation, context switches	
iostat -x 1	Disk I/O bottlenecks	Look for await, %util > 90%	
pidstat -p <pid> -u -r -w -d</pid>	Per-process CPU, mem, I/O	High %usr or minflt/majflt	
free -m	Memory, swap usage	Used/free/buffered swap	
dstat -tcmnd	Realtime combined stats	See CPU, mem, disk, network	
netstat -anp	Open ports and connections	Detect socket exhaustion	
lsof -p <pid></pid>	Open files/sockets	Detect FD leaks	
`ps -eLf	grep`	Thread count	
strace -p <pid></pid>	Syscall-level trace	Trace I/O, read/write, sleep	
perf top / perf record	CPU profiling	Hotspots at syscall/function level	
sar -u, sar -r, sar -B	Historic CPU/mem	Show trends in CPU, paging	
journalctl -u <service></service>	Systemd app logs	Debug startup/shutdown failure	

Q Log Patterns That Indicate Bottlenecks

Pattern	Where to Find It	Implication
OutOfMemoryError, Killed process	gc.log, kern.log, app logs	JVM killed due to memory limits
Segmentation fault, core dumped	syslog, dmesg	Native crash, bad C extension
blocked for more than 120 seconds	dmesg, kern.log	I/O lockup, filesystem stall
oom-killer events	dmesg, /var/log/kern.log	Memory exhaustion, swap mismanagement
connection reset by peer, broken pipe	App logs, network layer	Unhandled exception, retry storm
resource temporarily unavailable, EMFILE	dmesg, app logs	File descriptor exhaustion
too many open files	ulimit, Isof	Leaky sockets, handle leaks
permission denied, SELinux audit	audit.log, secure	Access issues

slow query, timeout	App/DB logs	Downstream slowness, API/DB
REJECT, DROP	iptables, auditd	Firewall or security blocking traffic

How to Approach a Unix-Level RCA (Root Cause Analysis)

1. Capture Key Metrics

uptime

vmstat 15

iostat -x 15

free -m

ps aux --sort=-%mem | head

2. Check for Kernel-Level Clues

dmesg | egrep -i 'error | warn | oom | blocked | segfault' journalctl -xe

3. Zoom into Specific Process

pidstat -p <pid> -u -r -w -d 1 strace -tt -T -p <pid> -o strace.log Isof -p <pid> | wc -l

4. Check Network/Port Exhaustion

netstat -an | grep ESTABLISHED | wc -l ss -s

5. Correlate with App Logs

grep -i 'exception' /var/log/<yourapp>/app.log grep -i 'ERROR' /var/log/<yourapp>/*

K Tools for Advanced Log & Trace Analysis

Tool	Usage	<mark>ldeal For</mark>
Logrotate	Manage log sizes & rotation	Prevent disk fill-up
logwatch, goaccess, Inav	Log summarization	Quick daily summaries

ELK Stack (Elasticsearch, Logstash, Kibana)	Centralized logging	Production log analysis
Fluentd, Filebeat	Shipping logs	K8s + cloud-native logs
bpftrace, eBPF	Kernel-level tracing	syscall/network profiling
perf, ftrace	High-res CPU tracing	Profiling Java/C binaries
systemtap, bcc-tools	Dynamic kernel probes	Low-level debugging
atop	Historic resource tracking	Find slow trends
Glances	All-in-one resource monitor	CLI + Web dashboard
Loggly, Datadog Logs, CloudWatch Logs	Cloud log aggregation	AWS-specific logging

Correlating UNIX Logs with Java Performance

Log/Event	Java Symptom	Unix Indicator
High GC time	CPU spikes, pause times	top, jstat, jstack, GC logs
Thread lock	High latency, timeouts	jstack, ps -eLf, strace shows blocking
I/O bottlenecks	API lag, DB slowness	iostat, dmesg, lsof
Memory leak	OOM errors, swap	dmesg, free -m, heap dump
File handle leak	Too many open files	Isof, ulimit -n exhaustion
Network stalls	Connection timeouts	netstat, tcpdump, ss

Recommended Aliases for Daily Use

```
alias top10cpu='ps aux --sort=-%cpu | head -n 11'
alias top10mem='ps aux --sort=-%mem | head -n 11'
alias threads='ps -eLf | wc -l'
alias fdcount='lsof -p $1 | wc -l'
alias memlog='free -h && vmstat 15'
alias cpuload='mpstat 15'
```

Heap Dump Analysis – The Ultimate Guide

What is a Heap Dump?

A Heap Dump is a binary snapshot of the JVM heap memory, containing:

- All live objects (and optionally GC roots, class metadata, thread stacks)
- Object types, counts, sizes, references

✓ Use Cases:

- Memory leaks
- OOM errors (java.lang.OutOfMemoryError: Java heap space)
- Cache growth analysis
- Object retention patterns
- ClassLoader leaks

How to Capture a Heap Dump

Method	Command/Tool	Notes Notes
On-Demand (Linux)	jmap -dump:format=b,file=heap.hprof <pid></pid>	May cause pause; works in prod if stable
Signal Trigger	kill -3 <pid> (for thread dump) + jcmd <pid> GC.heap_dump <file></file></pid></pid>	Safer than jmap in prod
At OOM	-XX:+HeapDumpOnOutOfMemoryError - XX:HeapDumpPath=/path/to/dump	Automatically on OOM error
JConsole/VisualVM	UI-based trigger	Easy for dev/stage; not ideal for prod
Lambda SnapStart	Part of snapshot; no manual dump	Use AWS Lambda tuning
EKS/ECS (Java in container)	docker exec <container> jcmd</container>	Heap dump inside container
Kubernetes	kubectl exec <pod> jcmd <pid> GC.heap_dump /tmp/heap.hprof</pid></pod>	For Java apps in K8s

🖺 What to Look For in a Heap Dump

Symptom Peap Dump Clue		Possible Root Cause & Fix
Memory Leak (High	Dominating classes in OldGenLarge	Unreleased references; caches without
OldGen Usage)	retained size	eviction; listeners
Unbounded Cache Growth	One or few classes with millions of instances	Map/Set/ConcurrentHashMap growth
ClassLoader Leak	Multiple ClassLoader instances holding references	Web apps, redeploys without cleanup
ThreadLocal Leaks	ThreadLocalMap holding objects	ThreadLocals not removed
Large Collections	Collections (List, Set, Map) with many elements	Poor batching, unbounded queues
Finalizer Retention	Objects in Finalizer queue	Overuse of finalize()
Duplicate Strings	Many identical String instances	String duplication
Excessive Buffers	Large arrays/byte[] objects	I/O, buffer leaks
Socket/Connection Leaks	Socket objects retained	Missing close() calls
Anonymous Inner Class Retention	Inner classes referencing outer	Large object graphs
Retained by GC Roots	Object graph leading to GC Roots (static fields, threads)	Strong references to static fields

Heap Dump Analysis Workflow (Battle-Tested)

- 1 Capture heap dump during OOM or memory spike.
- 2 Open with analysis tools:
 - **%** Eclipse Memory Analyzer (MAT)
 - X VisualVM
 - **K** GCEasy.io (for quick insights)
- 3 Run Leak Suspect Report (MAT) → Focus on Top Dominator Tree.
- 4 Check **Retained Size** vs. **Shallow Size**.
- 5 Find paths to GC Roots → Who's holding the object?
- 6 Investigate suspicious classes (e.g., java.util.HashMap, org.apache.catalina.loader.WebappClassLoaderBase).
- 7 Correlate with app code, caches, pools, and external systems.
- \blacksquare Test fix in dev/stage \rightarrow Recapture dump under load \rightarrow Compare results.

iii Heap Dump Tools

<mark>Tool</mark>	Strengths	<mark>ldeal Use</mark>
Eclipse Memory Analyzer (MAT)	Open-source, dominator tree, leak suspect report	Detailed offline analysis
VisualVM	Lightweight, live heap view, histograms	Quick checks, dev environments
GCEasy.io	Upload heap dump for instant report	Cloud-native insights
HeapHero.io	Reports + charts, classloader leaks	Cloud dumps
JProfiler	Heap + CPU profiling	Paid; great for dev/stage
YourKit	Advanced leak detection, object inspections	Paid; deep dives
JDK Mission Control (JMC)	Heap + thread + GC correlation	JFR integration

I JVM Heap Structure Cheat Sheet

<mark>Space</mark>	What It Holds	Flags to Tune
Eden	New objects, short-lived	-Xmn, G1 -XX:G1NewSizePercent
Survivor	Objects promoted between GCs	-XX:SurvivorRatio
OldGen	Long-lived objects	-Xmx, G1 -XX:InitiatingHeapOccupancyPercent
Metaspace	Class metadata	-XX:MaxMetaspaceSize
Code Cache	JIT-compiled code	-XX:ReservedCodeCacheSize

Advanced Heap Tuning Flags

Flag	Purpose
-XX:+HeapDumpOnOutOfMemoryError	Dump heap on OOM
-XX:HeapDumpPath=/path	Set heap dump location
-XX:+UseStringDeduplication	Reduce duplicate String instances (G1)
-XX:+PrintClassHistogram	Get class instance histogram
-XX:+UnlockDiagnosticVMOptions - XX:+PrintGCApplicationStoppedTime	Correlate GC with app pause time

LoadRunner Graph Correlation Matrix (Performance Analysis View)

Ø Primary Graph	© Correlates With	What It Tells You	🎇 Typical Root Cause
Average Response Time	- Throughput- Hits/sec- Errors/sec- CPU/Memory	If response time rises while throughput drops, likely server bottleneck	CPU overload, DB latency, GC pause, resource starvation
Throughput (Bytes/sec)	- Response Time- Hits/sec- Network I/O	High throughput with low response time = healthy. If throughput drops with rising response time, backend stress	Server under stress, DB wait, app server queueing
Hits/sec	- Concurrent Users- Throughput- CPU- TPS	Drop in hits/sec despite constant users may show errors, delays, queueing	Thread pool saturation, backend failure, bottlenecks
Transactions per Second (TPS)	- Response Time- Errors/sec- Backend logs	Stable TPS = stable system. Dips with response time spikes = app/DB delay	Slow SQL, code lock, thread pool exhaustion
Error Rate	- Response Time- TPS- Backend logs	Spikes in errors with stable load = functional defect or resource outage	Bad code path, 500 errors, broken dependency
Concurrent Vusers	- Hits/sec- TPS- Response Time	Constant vusers with rising response time = scalability bottleneck	Saturated queues, DB contention, I/O wait
Latency (Network)	- Response Time- DNS Time- TCP Connect Time	Spikes in latency can inflate response time → network issues	DNS, TLS handshake, bandwidth congestion
Pages Downloaded/sec	- Hits/sec- Throughput- TPS	Drop may indicate resource wait or asset timeout	CDN delay, image load failure, web server queueing
CPU Utilization (Server)	- Response Time- TPS- GC Activity	Spikes in CPU with drop in TPS = CPU saturation	Tight loops, serialization, poor concurrency
Memory Usage (Server)	- Response Time- GC logs- Swap Activity	Memory spikes = leak, cache growth, bad object lifetime	JVM heap leak, unbounded cache, unclosed objects

	T	T	T
GC Activity / GC Time	- Response Time- CPU	High GC time = memory churn ,	Poor object lifecycle,
	Usage- OldGen usage	causes response latency	small heap, high
			promotion
I/O Wait / Disk Usage	- Response Time- TPS-	High I/O wait = disk or DB	Slow IOPS, bad indexes,
	DB Logs	bottleneck	verbose logging
Network Usage	- Throughput-	Drops or spikes suggest payload	Compression, packet
(bytes in/out)	Hits/sec- Socket time	issue, delay, or congestion	loss, slow client
DB CPU or Time/Call	- TPS- Response Time-	Correlates high app latency with DB	Slow queries, missing
(Oracle, MySQL)	GC or thread dump	hotspots	indexes, lock
			contention
Think Time	- Response Time- TPS	Affects user pacing. If you remove	App can't scale linearly
(Controller)		think time and response time	
		degrades, system is near saturation	$oldsymbol{O}^*$
Transactions	- TPS- Errors/sec-	Shows functional correctness. Drop	Broken path, network
Passed/Failed	Hits/sec	in passed + spike in failed =	break, invalid response
		functional issue	
Bandwidth Graphs	- Throughput- Page	Detects network cap, latency	Network saturation,
	load- Network usage	issues, or CDN delays	AWS egress limit
JVM/GC Logs	- LoadRunner	Align GC spikes with LoadRunner	GC tuning needed
(external)	Response Time-	response latency	
	Server CPU		
Application Logs	- Error Rate-	Correlates server-side exceptions	Code defects,
(errors)	Response Time- TPS	with client-side metrics	unhandled errors, API
,			downtime

Practical Use Cases for Correlation

Scenario 1: Response Time ↑ + Throughput ↓ + CPU 100%

Likely Cause: CPU bottleneck on app servers

Action: Profile hot code paths, tune thread pools, optimize GC

Scenario 2: TPS ↓ + Errors ↑ + DB CPU ↑

Likely Cause: DB saturation or slow SQL

Action: Tune queries, check AWR, index, scale DB

⊘ Scenario 3: Hits/sec Stable + Response Time ↑ + GC Time ↑

Likely Cause: GC thrashing or memory leak

Action: Heap dump, GC log analysis, tune heap size/GC algorithm

Scenario 4: TPS ↓ + Hits/sec ↓ + Response Time Stable

Likely Cause: Load Generator saturation or network issue **Action**: Check LG CPU/mem; monitor network; check LG logs

Load Balancer Performance Bottlenecks & Correlation – Deep Dive

🗱 Load Balancer Metrics & What They Indicate

Metric 	What It Tells You	Symptoms	Root Cause	Action Plan
Request Count / New	Incoming request	Sudden spikes =	DDoS, bot traffic,	Rate limit, WAF rules,
Connections	volume	traffic surge	misbehaving clients	scale LB
Target Response Time	Time backend	High = slow	Backend bottleneck,	Profile app, DB, code
(ALB/NLB)	takes to respond	backend	DB latency	hotspots
Target Connection	Failed TCP/SSL	Spikes = backend	Backend app down,	Check backend
Errors	connections	overload	port closed, listener	health, port status
			misconfig	
HTTP 4XX / 5XX Rates	Client errors	Spikes in 5XX =	App crash, thread	App logs, thread
	(4XX) / Server	backend failure	pool starvation	dumps, GC logs
	errors (5XX)			
ELB 5XX Errors (502,	Gateway	504 = backend	Backend timeout,	Tune backend
504)	timeout, bad	timeout; 502 =	bad TLS certs	timeouts, fix TLS
	response	bad response		
Load Balancer Latency	Time LB takes to	Spikes during high	Queueing at LB,	Optimize target group
	route request	load	routing delays	health checks, enable
40				keep-alive
Healthy/Unhealthy	Backend instance	Drops = failing	App crash,	Check app logs,
Targets	health	instances	unhealthy	health check config
			endpoints	
Connection Reuse /	TCP connection	High new	Keep-alive disabled,	Enable keep-alive;
Recycle	handling	connections = no	misconfig	tune idle timeouts
		reuse		

SSL Negotiation Time	TLS handshake overhead	Long SSL times = latency	Large certs, no session resumption	Optimize certs, enable session reuse
Active Connections	Open TCP connections	High with low throughput = slow backend	Backend delay, DB lock, API slowness	Trace backend dependencies
HTTP Backend Errors (ALB)	App-side errors forwarded by LB	Spikes during traffic peaks	App error under load	Check app metrics, thread pool
Surge Queue (ALB	Requests queued	Growing =	Backend slow to	Scale backend,
Classic)	at LB	backend overload	respond	optimize app
Request Processing	Time LB spends	High = backend	DB/API bottleneck,	Profile backend
Time (ALB Classic)	per request	delay	app slowness	services
Dropped Connections	LB rejecting connections	Sudden spikes = overload	LB capacity, backend unresponsive	Scale LB, adjust timeouts

Load Balancer Bottleneck Patterns – Real-World Scenarios

Pattern	Symptom	Detection	Resolution
504 Gateway Timeout	User request hangs	ALB logs: 504s; X-Ray: backend	Tune backend timeouts,
	→ error	> timeout	increase app thread pools
502 Bad Gateway	App returns	ALB logs: 502s; App logs: 500	Check app code; TLS
	malformed response	errors	handshake issues
Latency Spikes at LB	Users report	LB metrics:	Backend saturation; thread
	slowness	TargetResponseTime 个	dump analysis; DB wait
High 4XX Errors	403/404/429 rates ↑	LB logs; WAF logs	Misconfigured routes, rate
			limiting, auth failures
Unhealthy Targets	Targets marked	LB health check logs	App crash, port closed,
	unhealthy		misconfig in health check
C'O			path
Connection Saturation	LB stops accepting	CloudWatch	Scale LB; adjust connection
	new connections	ActiveConnectionCount at	limits
		max	
TLS Handshake Delays	Users see slow initial	LB SSL metrics ↑	Use TLS 1.3; enable session
	load		tickets; optimize certs
Cross-AZ Traffic Cost	Unexpected AWS	VPC Flow Logs; LB metrics	Use same-AZ routing; check
	data transfer costs		target distribution

Keep-Alive Disabled	High new	CloudWatch: many	Enable keep-alive; tune
	connections; high	NewConnections	timeouts
	latency		
Load Balancer CPU High (Envoy/Nginx)	LB itself overloaded	Nginx/Envoy stats; OS CPU	Tune worker processes; offload TLS; distribute load
ALB/NLB Limits	New targets not	CloudWatch limits; AWS	Request limit increases;
Reached	registering	Service Quotas	refactor architecture

Correlating LB Metrics with Application & Infra

LB Metric	Correlates With	Root Cause Example
TargetResponseTime ↑	App Response Time 个, DB Latency 个	App slowness, DB locks
5XX Errors ↑	Thread Pool Saturation, GC Logs	App thread pool exhausted, GC pause
NewConnections ↑	DNS TTL, keep-alive reuse	TTL too low, keep-alive disabled
Latency ↑ + ActiveConns ↑	Backend CPU High	Backend CPU bottleneck under load
Unhealthy Targets ↑	App Crash Logs	App instance crash, port closed
TLS Negotiation Time ↑	Cert Size, Cipher Suites	Large certs, inefficient ciphers
Cross-AZ Traffic 个	VPC Flow Logs, Cost Explorer	ALB routing to different AZs
Surge Queue ↑	Backend Slow, Thread Dump	Backend app not scaling

Unix Internals: Swapping, Paging, Buffer I/O, and Context Switching

Core Concepts Cheat Sheet

Concept	Description	Impact if Bottlenecked
Swapping	Moving entire memory pages from RAM to disk (swap space)	Severe performance hit; processes freeze
Paging	Kernel moves memory pages between RAM and swap (demand paging)	Increased page faults; degraded performance
Buffer I/O	Kernel buffers disk I/O (page cache, dirty pages)	I/O latency spikes; high disk wait time

Context	CPU switches from one process/thread to another	High switches = CPU overhead,
Switching		cache misses

III Bottleneck Indicators and Metrics

Symptom	Metric/Command	Interpretation	Root Cause
High Swap Usage	free -m, swapon -s, vmstat si/so	swap used > 0, si/so > 0 = active swapping	Memory pressure; OOM risk
High Page Faults	vmstat -s, sar -B	pgfault/s (minor), pgmajfault/s (major)	Excessive memory access; cache misses
Disk I/O Wait	iostat -x, vmstat wa	%util > 90%, await > 20ms, wa > 10%	Slow disk; excessive buffer flushes
High Context	vmstat cs, sar -w, pidstat	>10k/s (app-specific) = potential	Mutex locks, threads, I/O
Switch Rate	-w	lock contention	waits
Dirty Pages	/proc/meminfo, vmstat	Dirty pages growing; sync	Slow disk writeback; fsync
Accumulating		lagging	bottleneck
High r in vmstat	vmstat 1	r > CPU cores = CPU saturation	Process queue buildup; CPU bottleneck

Swapping Analysis

Check	Command	Bottleneck Sign
Swap Usage	free -m, swapon -s	Swap > 0 MB = memory pressure
Swap In/Out Rate	vmstat 1 (si, so)	Non-zero si/so = active swapping
Major Page Faults	sar -B	High pgmajfault/s = swap-in delays
dmesg Logs	`dmesg	grep -i oom`

Action:

- Tune app memory; reduce cache sizes; increase RAM.
- Disable swap (swapoff -a) for performance tests (careful!).

Paging Analysis

Check	Command	Bottleneck Sign
Page Fault Rate	vmstat -s, sar -B	High pgfault/s + pgmajfault/s
Minor vs Major	vmstat -s	High major = slow; minor = normal
Page Reclaims	vmstat -s	High reclaims = page pressure

Action:

- Optimize app memory locality.
- Avoid large random memory allocations.
- Reduce memory-hungry processes.

Buffer I/O Analysis

Check	Command	Bottleneck Sign
Dirty Pages	`cat /proc/meminfo	grep Dirty`
Disk Wait	iostat -x 1, vmstat wa	High %util, await, wa
Writeback Latency	dstat -d, sar -b	High blocks out; delayed I/O

Action:

- Tune vm.dirty_ratio, vm.dirty_background_ratio:
- sysctl -w vm.dirty_ratio=10
- sysctl -w vm.dirty_background_ratio=5
- Use faster disks (SSD/NVMe); distribute I/O load.

Context Switch Analysis

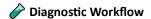
Check	Command	Bottleneck Sign	
Context Switch Rate	vmstat cs, pidstat -w, sar -w	>10k/s sustained = high overhead	
Thread Count	`ps -eLf	wc -l`	
Lock Contention	perf top, strace -p <pid></pid>	Threads in futex, epoll_wait	

Action:

- Reduce threads; pool threads effectively.
- Refactor synchronized code; reduce lock granularity.
- Avoid excessive I/O-bound threads.

Correlation Matrix

Symptom	Swapping	Paging	Buffer I/O	Context Switch
App Latency 个	High swap usage	High major faults	High dirty pages, wa	High context switches
Disk I/O Wait 个	Low RAM = swap-in	High minor faults	High await, %util	Threads waiting on I/O
CPU Spikes	Swap-in CPU	Page fault handling	I/O interrupts	High task switches
OOM Kills	Yes	Yes	No	Indirect (lock wait stalls)



1 Check RAM:

free -m

2 Monitor Swapping:

vmstat 1

3 Page Faults:

sar -B 1 5

4 Disk I/O Wait:

iostat -x 1

5 Context Switches:

pidstat -w 1

6 Correlate with App Logs, GC, and Thread Dumps

% Advanced Tools

Tool	Purpose	
perf top, perf record	Kernel-level hotspots	
bpftrace, bcc-tools	Trace syscalls, I/O wait	
strace -p <pid></pid>	Syscall-level blocking	
iotop, dstat	Per-process disk I/O	
sar, atop, glances	Combined metrics overview	