

Comprehensive Log Analysis Guide

1. General Log Analysis

Using `grep`

1. Find lines with specific keywords (case-insensitive):

```
grep -i "error" application.log
```

- **What It Does:** Searches for lines containing the keyword "error" (case-insensitive).
- **Example Use Case:** Identify all error occurrences in a log file.

2. Search for multiple keywords:

```
grep -iE "error|failed|exception" application.log
```

- **What It Does:** Matches lines containing "error", "failed", or "exception".
- **Example Use Case:** Debugging application logs for different error patterns.

3. Count occurrences of a keyword:

```
grep -i "timeout" application.log | wc -l
```

- **What It Does:** Counts all lines containing "timeout".
- **Example Output:** 45 (indicating 45 timeout errors).

4. Find lines with context (before/after matches):

```
grep -i "error" application.log -A 5 -B 3
```

- **What It Does:** Displays 5 lines after and 3 lines before each match.
- **Example Use Case:** Analyze logs surrounding errors for better context.

5. Exclude lines containing specific keywords:

```
grep -i "error" application.log | grep -vi "debug"
```

- **What It Does:** Filters out lines containing "debug" from the error matches.

6. Search for lines matching specific patterns (e.g., IP addresses):

```
grep -oE "([0-9]{1,3}\.){3}[0-9]{1,3}" access.log
```

- **What It Does:** Extracts IPv4 addresses from the log file.

7. Monitor logs in real-time for specific issues:

```
tail -f system.log | grep --line-buffered -i "critical"
```

- **What It Does:** Streams new log entries and filters for "critical".

8. Search for lines within a specific date range:

```
grep -E "2025-01-05|2025-01-06" application.log | grep -i "error"
```

- **What It Does:** Finds errors logged on specific dates.
-

Using `awk`

1. Print specific fields from matched lines:

```
awk '/error/ {print $1, $2, $5}' application.log
```

- **What It Does:** Prints the first, second, and fifth fields from lines containing "error".

2. Summarize errors by hour:

```
awk '/error/ {split($2, time, ":"); hour[time[1]]++;} END {for (h in hour) print h, hour[h]}' application.log
```

- **What It Does:** Groups and counts errors by the hour.
- **Example Output:**

```
10 35
11 28
12 40
```

3. Summarize errors by type:

```
awk '/error/ {type[$3]++;} END {for (t in type) print t, type[t]}' application.log | sort -nr -k2
```

- **What It Does:** Groups and counts errors based on the third field.

4. Calculate average response times:

```
awk '{sum += $NF} END {print "Average response time:", sum/NR}' access.log
```

- **What It Does:** Computes the average response time, assuming the last field is the time.
-

Using sed

1. Highlight specific keywords:

```
sed 's/error/\x1b[31m&\x1b[0m/lg' application.log
```

- **What It Does:** Highlights "error" in red for improved readability.

2. Extract lines between specific markers:

```
sed -n '/START_ERROR/,/END_ERROR/p' application.log
```

- **What It Does:** Extracts all lines between START_ERROR and END_ERROR.

3. Remove sensitive information (e.g., email addresses):

```
sed -E 's/[a-zA-Z0-9._%+~]+@[a-zA-Z0-9.-]+\.[a-zA-Z]{2,}/[REDACTED]/g' logs.log
```

- **What It Does:** Masks email addresses in the logs.
-

Using head and tail

1. Display the first 20 lines of a log file:

```
head -20 application.log
```

- **What It Does:** Shows the first 20 lines.

2. Display the last 50 lines containing "error":

```
grep -i "error" application.log | tail -50
```

- **What It Does:** Extracts the last 50 matches for "error".

3. Monitor real-time logs:

```
tail -f application.log
```

- **What It Does:** Streams new log entries in real-time.
-

2. Database Log Analysis

Using `grep`

1. Find Oracle errors (e.g., ORA- codes):

```
grep "ORA-" oracle.log
```

- **What It Does:** Matches lines with Oracle error codes.

2. Search for slow queries:

```
grep "execution time: [5-9][0-9]*ms" db.log
```

- **What It Does:** Matches SQL queries taking over 5 seconds.
-

Using `awk`

1. Summarize Oracle errors by code:

```
awk '/ORA-/ {split($0, a, ":"); print a[1]} oracle.log | sort | uniq -c | sort -nr
```

- **What It Does:** Groups and counts occurrences of each Oracle error code.

2. Identify largest table space usage:

```
awk '/tablespace/ {print $2, $5} oracle.log | sort -k2 -nr | head -5
```

- **What It Does:** Lists the top 5 table spaces by usage.
-

3. JVM Log Analysis

Using `grep`

1. Find memory-related errors:

```
grep -i "OutOfMemoryError" jvm.log
```

- **What It Does:** Matches memory-related errors in JVM logs.

2. Search for specific exceptions:

```
grep -A 10 -B 5 "NullPointerException" jvm.log
```

- **What It Does:** Displays 10 lines after and 5 lines before NullPointerException.
-

4. Kubernetes Log Analysis

Using `grep`

1. Find pods in restart loops:

```
grep -i "CrashLoopBackOff" kube.log
```

- **What It Does:** Matches pods stuck in CrashLoopBackOff.

2. Detect out-of-memory kills:

```
grep -i "OOMKilled" kube.log
```

- **What It Does:** Finds containers terminated due to memory issues.
-

Using `awk`

1. Count restarts for each pod:

```
awk '/Restarted/ {pod=$2; restarts[pod]++} END {for (p in restarts) print restarts[p], p}' kube.log | sort -nr | head -10
```

- **What It Does:** Lists the top 10 pods with the highest restarts.
-

5. AWS Log Analysis

Using `grep`

1. Find Lambda errors:

```
grep -i "error" aws-lambda.log | grep -i "function"
```

- **What It Does:** Matches errors related to AWS Lambda functions.

2. Search for API Gateway 5xx errors:

```
grep '"status":5[0-9][0-9]' api-gateway.log
```

- **What It Does:** Matches HTTP 5xx status codes.

6. Automating Log Analysis with Scripts

Daily Error Summary Script

```
#!/bin/bash

LOGFILE="application.log"
OUTPUT="error_summary_$(date +%F).txt"

echo "Error Summary for $(date)" > "$OUTPUT"
echo "-----" >> "$OUTPUT"

# Count total errors
echo "Total Errors:" >> "$OUTPUT"
grep -i "error" "$LOGFILE" | wc -l >> "$OUTPUT"

# Summarize errors by type
echo "Error Types:" >> "$OUTPUT"
grep -i "error" "$LOGFILE" | awk '{print $3}' | sort | uniq -c | sort -nr >> "$OUTPUT"

# Show the top 10 error lines
echo "Top 10 Errors:" >> "$OUTPUT"
grep -i "error" "$LOGFILE" | head -10 >> "$OUTPUT"

echo "Summary generated successfully at $(date)." >> "$OUTPUT"
```

7. Combining Tools for Advanced Use Cases

Real-Time Monitoring with Highlighting

```
tail -f system.log | grep --line-buffered -i "critical" | sed 's/critical/\x1b[31m&\x1b[0m/'
```

- **What It Does:** Monitors new log entries, highlights "critical" in red, and displays them in real-time.
-