

# Kibana Setup, Prerequisites, Workflow, and Navigation

## Flow of Log Aggregation and Dashboards

**Kibana** is an open-source visualization tool that works with Elasticsearch to provide real-time analysis, visualization, and monitoring of logs and structured data. It is primarily used in observability, security, and business intelligence.

---

### 1. Prerequisites for Kibana Setup

Before setting up Kibana, ensure the following prerequisites:

#### 1.1 Infrastructure Requirements

- **OS:** Kibana can run on Linux, Windows, or macOS.
- **RAM:** At least **4GB of RAM** (recommended: 8GB+ for production).
- **CPU:** Minimum **2 vCPUs** (recommended: 4+ vCPUs).
- **Storage:** SSD storage with high IOPS is recommended.
- **Networking:** Ensure Kibana can communicate with Elasticsearch on port **9200**.

#### 1.2 Software Requirements

- **Elasticsearch:** Kibana is tightly integrated with Elasticsearch and requires a running Elasticsearch cluster.
- **Java** (for Elasticsearch, if applicable).
- **Node.js** (optional for development).
- **Nginx or Apache** (optional for reverse proxy and authentication).
- **Logstash / Beats** (optional for log aggregation).

#### 1.3 Installation Steps

##### 1.3.1 Install Kibana

1. **Download Kibana** from the [Elastic website](#).
2. Extract the package:

```
tar -xzf kibana-<version>-linux-x86_64.tar.gz
```

```
cd kibana-<version>-linux-x86_64
```

3. Edit the configuration file (kibana.yml):

```
server.port: 5601
```

server.host: "0.0.0.0"

elasticsearch.hosts: ["http://localhost:9200"]

4. Start Kibana:

./bin/kibana

5. Open **http://localhost:5601/** in a web browser.

### 1.3.2 Configure Elasticsearch for Kibana

- Ensure that the Elasticsearch cluster is up and running.
- Update the Elasticsearch configuration (elasticsearch.yml):

cluster.name: my-cluster

node.name: node-1

network.host: 0.0.0.0

http.port: 9200

---

## 2. Workflow and Navigation Flow of Log Aggregation and Dashboards

### 2.1 Log Aggregation Process

#### 1. Log Collection:

- Use **Filebeat**, **Logstash**, or **Fluentd** to collect logs.
- Example: Filebeat configuration (filebeat.yml):

filebeat.inputs:

- type: log

paths:

- /var/log/nginx/access.log

output.elasticsearch:

hosts: ["http://localhost:9200"]

- Start Filebeat:

sudo service filebeat start

#### 2. Log Ingestion and Parsing:

- Logstash pipelines can filter, transform, and enrich logs before sending them to Elasticsearch.

- Example Logstash configuration (logstash.conf):

```
input {
  beats {
    port => 5044
  }
}

filter {
  grok {
    match => { "message" => "%{COMMONAPACHELOG}" }
  }
}

output {
  elasticsearch {
    hosts => ["http://localhost:9200"]
  }
}
```

### 3. Indexing in Elasticsearch:

- Logs are stored in indices with a schema.
- Example index creation:

```
curl -X PUT "localhost:9200/logs-2025.02.06"
```

### 4. Visualization in Kibana:

- Create **Data Views** (formerly known as Index Patterns) in Kibana.
- Build visualizations and dashboards.

---

## 3. Key Components of Kibana

### 3.1 Data View (Formerly Index Patterns)

- Allows Kibana to query specific indices in Elasticsearch.
- Example:
  - Go to **Management > Stack Management > Data Views**.

- Click **Create Data View**.
- Enter logs-\* as the index pattern.

### 3.2 Discover

- Explore raw logs with filters.
- Useful for debugging performance bottlenecks.

### 3.3 Visualizations

- Types of visualizations:
  - **Line charts** (for trends)
  - **Bar charts** (categorical data)
  - **Pie charts** (proportions)
  - **Maps** (geospatial data)
  - **Metric visualizations** (KPIs)

### 3.4 Dashboards

- Combine multiple visualizations into a single view.
- Example: Create a dashboard for **application latency monitoring**.
  1. Go to **Dashboards > Create Dashboard**.
  2. Add visualizations like:
    - **Average response time** (line chart).
    - **CPU/memory usage trends** (area chart).
    - **Error rate** (bar chart).

### 3.5 Alerts & Reporting

- Create alerts for anomalies (e.g., high latency).
- Example: If response time exceeds 500ms, send an alert to Slack.

---

## 4. Detecting Performance Issues in Kibana

Kibana can be used to monitor **application performance, infrastructure health, and bottlenecks** using various techniques.

### 4.1 Key Metrics to Track

- **Response time** (avg(response\_time))

- **Error rate** (count(status\_code: 500))
- **CPU and memory utilization** (avg(cpu\_usage))
- **Database query performance** (avg(query\_execution\_time))
- **Elasticsearch query latency** (search.latency)

## 4.2 Techniques for Performance Issue Detection

### 4.2.1 Slow Queries Analysis

#### 1. Enable Elasticsearch Slow Query Logs

- Update elasticsearch.yml:
 

```
index.search.slowlog.threshold.query.warn: 2s
index.search.slowlog.threshold.fetch.warn: 1s
```
- Query logs in Kibana:
 

```
{
  "query": {
    "match": {
      "message": "took more than 2s"
    }
  }
}
```

#### 2. Create Kibana Visualization

- Use "Query Time vs Request Count" visualization.
- Identify peak query times.

### 4.2.2 Detecting JVM GC Issues

#### 1. Ship GC logs to Elasticsearch:

- Use **Filebeat or Logstash**.
- Example Logstash filter:
 

```
filter {
  grok {
    match => { "message" => "%{NUMBER:gc_time}" }
  }
}
```

```
}
```

## 2. Analyze in Kibana:

- Create a dashboard for gc\_time trends.
- Set alerts for excessive GC pauses.

### 4.2.3 Memory & CPU Bottlenecks

- Use **Metricbeat** to collect CPU and memory metrics.
- Example metricbeat.yml:

```
metricbeat.modules:  
  - module: system  
  
metricsets:  
  - cpu  
  - memory  
  
period: 10s  
  
output.elasticsearch:  
  hosts: ["http://localhost:9200"]
```

- Kibana dashboard can show:
  - **CPU usage spikes.**
  - **High memory consumption.**
  - **Correlate with application logs.**

### 4.2.4 Analyzing Load Balancer and API Gateway Logs

- If using **AWS ALB logs**, ship them via Logstash.
- Example filter:

```
filter {  
  grok {  
  
    match => { "message" => "%{TIMESTAMP_ISO8601:timestamp}  
%{NUMBER:response_time}" }  
  
  }  
}
```

- Create a Kibana dashboard to analyze:

- **Response time trends.**
  - **High error rate periods.**
  - **Spikes in load balancer latency.**
-

# Detecting Performance Issues in Kibana – A Technical Deep Dive

It's absolutely possible to detect and troubleshoot performance issues using **Kibana**, provided it is integrated with **Elasticsearch**, **Logstash**, **Metricbeat**, and **APM (Application Performance Monitoring)**. Kibana enables **real-time monitoring**, **log aggregation**, **data visualization**, and **alerting** to identify bottlenecks in application and infrastructure performance.

---

## 1. Key Performance Issues That Kibana Can Detect

Kibana, in combination with Elasticsearch and Beats, can help identify the following **performance bottlenecks**:

### 1. Application Performance Issues

- High response times for API requests.
- Increased error rates (HTTP 500, 503, etc.).
- Slow database queries affecting user experience.
- Garbage Collection (GC) pauses affecting Java applications.

### 2. Infrastructure Issues

- High CPU and memory usage on application servers.
- Network latency and packet loss in Kubernetes clusters.
- Load balancer latency affecting request routing.

### 3. Database Performance Issues

- Slow SQL queries with high execution times.
- Increased database connection wait times.
- High database CPU usage.

### 4. Elasticsearch Performance Issues

- High query response time in Elasticsearch.
  - Slow indexing speed due to large ingestion rates.
  - High heap usage leading to frequent GC pauses.
-



## 2. Key Techniques for Detecting Performance Issues in Kibana

To detect performance issues using Kibana, we follow **multiple techniques**, including **real-time monitoring, anomaly detection, time-series analysis, and alerting**.

### 2.1 Log Analysis for Application Performance Issues

By ingesting application logs into **Elasticsearch**, we can detect slow response times, errors, and exceptions.

#### Step 1: Enable Application Logging

If using Java Spring Boot, configure logging with ELK:

logging:

level:

root: INFO

org.springframework.web: DEBUG

org.hibernate: ERROR

Ensure logs are shipped via **Filebeat**:

filebeat.inputs:

- type: log

paths:

- "/var/log/app.log"

output.elasticsearch:

hosts: ["http://localhost:9200"]

Start Filebeat:

sudo service filebeat start

#### Step 2: Create a Data View in Kibana

- Navigate to **Stack Management > Data Views > Create Data View**.
- Select app-logs-\* as the index pattern.

#### Step 3: Build a Dashboard

- Create a **line chart visualization** for:
  - **Average Response Time** (avg(response\_time))
  - **Error Rate** (count(status\_code: 500))

- Use the **"Discover" tab** to filter logs:

```
{  
  "query": {  
    "match": {  
      "status_code": "500"  
    }  
  }  
}
```

- Set alerts if response time exceeds a threshold (e.g., **>500ms**).

---

## 2.2 Monitoring System Metrics (CPU, Memory, Disk, Network)

System-level performance issues like **high CPU, memory leaks, and disk bottlenecks** can be monitored using **Metricbeat**.

### Step 1: Configure Metricbeat

Install and configure **Metricbeat**:

```
sudo apt-get install metricbeat
```

Edit the metricbeat.yml file:

```
metricbeat.modules:
```

```
- module: system
```

```
  metricsets: ["cpu", "memory", "network"]
```

```
  period: 10s
```

```
output.elasticsearch:
```

```
  hosts: ["http://localhost:9200"]
```

Start Metricbeat:

```
sudo service metricbeat start
```

### Step 2: Build a Dashboard

- Use **"Lens" visualization** in Kibana to create:
  - **CPU Usage (%) Over Time** (system.cpu.total.pct)
  - **Memory Consumption Trends** (system.memory.used.pct)

- **Disk I/O Latency**
- **Network Packet Drops**

### Step 3: Set Alerts

- **Trigger alerts** when CPU usage > **85%** for 5 minutes.
- **Trigger alerts** if memory utilization crosses **90%**.

## 2.3 Slow Query Analysis for Database Performance

Database slow queries often impact overall system performance.

### Step 1: Enable Slow Query Logs

For **MySQL**, modify my.cnf:

```
slow_query_log = 1
long_query_time = 2
log_output = FILE
slow_query_log_file = /var/log/mysql-slow.log
```

For **PostgreSQL**, modify postgresql.conf:

```
log_min_duration_statement = 2000
log_statement = 'all'
```

### Step 2: Ship Logs to Elasticsearch

Modify **Filebeat configuration** (filebeat.yml):

```
filebeat.inputs:
  - type: log
    paths:
      - "/var/log/mysql-slow.log"
```

output.elasticsearch:

```
hosts: ["http://localhost:9200"]
```

Start Filebeat:

```
sudo service filebeat restart
```

### Step 3: Create a Kibana Dashboard

- **Visualize slow queries by execution time.**

- Query example:

```
{
  "query": {
    "range": {
      "query_time": {
        "gte": "2s"
      }
    }
  }
}
```

#### Step 4: Set Alerts

- If **query execution time** exceeds **5s**, trigger an alert.
- If **DB connections exceed a threshold**, send a Slack alert.

## 2.4 Elasticsearch Performance Monitoring

Since Kibana relies on **Elasticsearch**, monitoring Elasticsearch itself is critical.

### Step 1: Enable Slow Query Logs

Modify elasticsearch.yml:

```
index.search.slowlog.threshold.query.warn: 2s
index.search.slowlog.threshold.fetch.warn: 1s
```

### Step 2: Query Elasticsearch Metrics in Kibana

Query Elasticsearch for slow queries:

```
{
  "query": {
    "match": {
      "message": "took more than 2s"
    }
  }
}
```

### Step 3: Set Up an Elasticsearch Performance Dashboard

- Query response time trend (search.latency)
- Slow indexing rate (indexing.latency)
- JVM heap usage (jvm.mem.heap\_used\_percent)

### Step 4: Detect High Heap Usage

If JVM Heap usage > **75%**, potential **GC issues**. Create alerts to restart Elasticsearch nodes if necessary.

---

## 2.5 APM-Based Application Performance Monitoring

Using **Elastic APM**, we can trace application transactions and detect slow endpoints.

### Step 1: Install APM Server

```
sudo apt-get install apm-server
```

Modify apm-server.yml:

```
apm-server:
```

```
host: "0.0.0.0:8200"
```

```
output.elasticsearch:
```

```
hosts: ["http://localhost:9200"]
```

Start APM:

```
sudo service apm-server start
```

### Step 2: Integrate with Java Application

Add the **Elastic APM agent**:

```
<dependency>  
<groupId>co.elastic.apm</groupId>  
<artifactId>apm-agent-attach</artifactId>  
<version>1.27.0</version>  
</dependency>
```

Start the application with APM:

```
java -javaagent:/path/to/elastic-apm-agent.jar -  
Delastic.apm.server_urls=http://localhost:8200 -jar myapp.jar
```

### Step 3: Create APM Dashboards

- Track slow transactions.
  - Monitor database query execution time.
  - Identify high-latency APIs.
- 

### 3. Conclusion

- ✓ Kibana can detect application, system, database, and Elasticsearch performance issues.
  - ✓ Using Filebeat, Metricbeat, and APM, you can ingest logs and metrics into Elasticsearch for analysis.
  - ✓ Dashboards and alerts help detect bottlenecks in real time.
  - ✓ Anomaly detection can identify performance degradations before they impact users.
- 

### Case Study 1: Slow Application Response Times Due to High Load

#### Scenario:

A banking application was experiencing **intermittent slow response times**, especially during peak hours. The latency of some API endpoints exceeded **5 seconds**, causing timeouts.

#### Step 1: Data Collection

- **Logs** from the application were collected using **Filebeat**.
- **Metrics** (CPU, memory, thread pool) were collected using **Metricbeat**.
- **Transaction traces** were collected using **Elastic APM**.

#### Step 2: Kibana Dashboard Setup

A **custom dashboard** was created in Kibana to monitor:

- **Average API Response Time (ms)**
- **Error Rate (%)**
- **JVM Heap Usage (%)**
- **CPU Utilization (%)**

### Query to Identify Slow API Calls:

```
{
  "query": {
    "range": {
      "response_time": {
        "gte": "5000"
      }
    }
  }
}
```

### Visualization: API Response Time Over Time

- A **line chart** plotted API response times.
- A **bar chart** visualized slow API calls grouped by endpoint.

### Step 3: Root Cause Analysis

#### 1. High CPU Usage:

- Kibana showed **CPU spikes above 90%** during peak hours.
- JVM **Garbage Collection (GC) logs** showed **long GC pauses (avg. 2s)**.

#### 2. Thread Pool Exhaustion:

- **Thread dump analysis** indicated that the application **ran out of worker threads**, causing request queueing.
- Too many active threads in `java.util.concurrent.ThreadPoolExecutor`.

#### 3. Database Bottleneck:

- Database logs (shipped via **Filebeat**) showed **slow SQL queries** running over **3 seconds**.
- Query example:

```
SELECT COUNT(*) FROM transactions WHERE user_id = ?;
```

### Step 4: Resolution & Optimization

✓ **JVM tuning:** Increased heap memory allocation and enabled **G1GC** for better memory management.

- ✓ **Thread pool tuning:** Increased worker thread count and optimized connection pool settings.
- ✓ **Database indexing:** Added indexes to frequently queried columns in the transactions table.

#### Alerting Configuration:

- Trigger alert **if API response time exceeds 2s for 5 minutes.**
- Send Slack notifications using:

```
{  
  "actions": {  
    "slack_notification": {  
      "message": "High API latency detected (>2s). Check JVM and DB logs."  
    }  
  }  
}
```

---

#### Case Study 2: Elasticsearch Slow Query Performance

##### Scenario:

An **e-commerce platform** using **Elasticsearch** for product searches reported **slow query performance**, affecting customer experience.

##### Step 1: Enabling Slow Query Logs

Elasticsearch **slow query logging** was enabled:

```
index.search.slowlog.threshold.query.warn: "2s"  
index.search.slowlog.threshold.fetch.warn: "1s"
```

##### Step 2: Kibana Dashboard Setup

A **custom dashboard** in Kibana monitored:

- **Query Execution Time**
- **Search Latency Over Time**
- **Indexing Performance**

##### Query to Find Slow Queries

```
{  
  "query": {
```

---



```

"range":{
  "took":{
    "gte": "2000"
  }
}
}
}
}
}

```

### Visualization: Query Latency Histogram

- A **histogram** grouped slow queries by execution time.
- A **heatmap** mapped query execution times against product categories.

### Step 3: Root Cause Analysis

#### 1. Large Result Sets:

- Some **search queries** returned **100,000+ results**, causing excessive data transfer.
- Example slow query:

```

{
  "query":{
    "match":{
      "product_description": "wireless headphones"
    }
  },
  "size": 100000
}

```

- Solution: **Limit size to 50** and implement **pagination**.

#### 2. Expensive Aggregations:

- Kibana logs showed **slow queries with aggregations**:

```

{
  "aggs":{
    "avg_price":{
      "avg":{

```

```

        "field": "price"
      }
    }
  }
}

```

- Solution: Precompute aggregation results in a **separate index**.

### 3. High JVM Heap Usage:

- Elasticsearch **JVM heap exceeded 80%**, leading to **frequent GC pauses**.
- Solution: **Increase heap size** (-Xms8g -Xmx8g) and **use circuit breakers**.

### Step 4: Resolution & Optimization

✓ **Query optimization:** Used match\_phrase instead of match for better relevancy.

✓ **Index sharding:** Increased the number of **shards from 3 to 6** to improve search distribution.

✓ **Elasticsearch caching:** Enabled request\_cache for repeated queries.

### Alerting Configuration:

- If **search latency exceeds 2s**, trigger an alert in Kibana.
- Log alert messages:

```

{
  "actions": {
    "email_notification": {
      "message": "Elasticsearch query latency exceeded 2s. Investigate!"
    }
  }
}

```

---

## Case Study 3: Kubernetes Network Latency Issues

### Scenario:

A **Kubernetes-based microservices system** was experiencing **intermittent high network latency**, causing API timeouts.

### Step 1: Collecting Logs and Metrics

- **Filebeat** collected application logs.

- **Metricbeat** collected CPU, memory, and **network statistics**.

## Step 2: Kibana Dashboard Setup

The following visualizations were created:

- **Latency Heatmap:** API response times across regions.
- **Pod-Level Network Traffic:** Network requests per Kubernetes pod.

## Query to Detect High Latency

```
{
  "query": {
    "range": {
      "latency": {
        "gte": "2000"
      }
    }
  }
}
```

## Step 3: Root Cause Analysis

### 1. Packet Drops in Kubernetes Pods

- Kibana logs showed **high packet drops** in specific nodes:

```
{
  "event.dataset": "system.network",
  "network.dropped.packets": ">1000"
}
```

### 2. AWS Load Balancer Latency

- Analyzed **AWS ALB logs** and found **p99 latency spikes > 3s**.

### 3. Intermittent DNS Resolution Delays

- CoreDNS logs showed **delays resolving service names**.

#### Step 4: Resolution & Optimization

- ✓ Increased Kubernetes pod replicas to distribute load.
- ✓ Enabled CoreDNS caching for faster DNS resolution.
- ✓ Optimized AWS ALB configurations by enabling HTTP/2 keep-alive.

#### Alerting Configuration:

- If API latency exceeds 1s, trigger an alert:

```
{  
  "actions": {  
    "slack_notification": {  
      "message": "High network latency detected! Check Kubernetes pods."  
    }  
  }  
}
```

---

# Hands-on Examples for Setting Up Alerts, Queries, and Visualizations in Kibana

Below are **step-by-step** examples to set up **queries, visualizations, and alerts** in **Kibana** for performance monitoring. These examples cover **application monitoring, database slow queries, Elasticsearch performance, and Kubernetes network latency**.

---

## 1. Setting Up Queries in Kibana (Discover Tab)

### Example 1: Querying Slow API Requests

If an API response time exceeds **500ms**, we need to track these slow responses.

#### Steps:

1. Open **Kibana > Discover**.
2. Select the appropriate **Data View** (formerly Index Pattern), e.g., `app-logs-*`.
3. Use the following **KQL (Kibana Query Language)** filter:

`response_time > 500`

4. Click **"Update"** to see results.
5. Save the query for future analysis.

### Equivalent Elasticsearch JSON Query

If querying Elasticsearch directly:

```
{
  "query": {
    "range": {
      "response_time": {
        "gte": 500
      }
    }
  }
}
```

## Example 2: Detecting Frequent HTTP 500 Errors

To find all logs where the application returns **HTTP 500** errors:

### KQL Query in Kibana Discover

status\_code: 500

### Equivalent Elasticsearch Query

```
{
  "query": {
    "match": {
      "status_code": "500"
    }
  }
}
```

---

## Example 3: Identifying Slow Database Queries

If SQL queries take more than **2 seconds**, use this query.

### KQL Query

query\_time > 2000

### Elasticsearch JSON Query

```
{
  "query": {
    "range": {
      "query_time": {
        "gte": 2000
      }
    }
  }
}
```

---

## 2. Creating Visualizations in Kibana

### Example 1: Visualizing API Response Time Trends

1. Go to Kibana → Visualize.
  2. Click **Create Visualization**.
  3. Select **Line Chart**.
  4. Choose the data source **app-logs\***.
  5. Set:
    - **Y-axis:** Average response\_time.
    - **X-axis:** Time (@timestamp).
  6. Click **Save** and add it to a **Dashboard**.
- 

### Example 2: Slow Query Distribution in Database

1. Go to Kibana → Visualize.
  2. Click **Create Visualization**.
  3. Choose **Bar Chart**.
  4. Select **db-logs\*** as the data source.
  5. Configure:
    - **Y-axis:** Count of logs.
    - **X-axis:** Terms aggregation on query\_time.
    - Add a filter: query\_time > 2000.
  6. Click **Save** and add it to a **Dashboard**.
- 

## 3. Creating Dashboards in Kibana

### Example: Application Performance Dashboard

1. Go to Kibana → Dashboard.
2. Click **Create New Dashboard**.
3. Click **Add Existing Visualizations**:
  - API Response Time Trend (Line Chart).
  - Slow Database Queries (Bar Chart).

- HTTP 500 Errors (Metric Count).
4. Save the dashboard as **"App Performance Dashboard"**.
- 

#### 4. Setting Up Alerts in Kibana

##### Example 1: Alert for High API Response Times

1. Go to Kibana → Stack Management → Rules & Connectors.
2. Click **Create Rule**.
3. Choose **Elasticsearch Query**.
4. Set:

- Name: **High API Latency Alert**
- Index: **app-logs-\***.
- Query:

```
{  
  "query": {  
    "range": {  
      "response_time": {  
        "gte": 1000  
      }  
    }  
  }  
}
```

5. Set a **Trigger Condition**:
  - If response\_time > 1000ms for 5 minutes.

6. **Choose Action**:

- Send an **email alert** or **Slack notification**.

- Message:

```
{  
  
  "message": "Warning! API response time exceeded 1s. Investigate possible performance  
issues."  
  
}
```



7. Save the rule.
- 

### Example 2: Alert for Elasticsearch Slow Queries

1. Go to Kibana → Stack Management → Rules & Connectors.
2. Click **Create Rule**.
3. Choose **Elasticsearch Query**.
4. Set:

- Name: **Slow Query Alert**
- Index: **es-logs-\***.
- Query:

```
{  
  "query": {  
    "range": {  
      "took": {  
        "gte": 2000  
      }  
    }  
  }  
}
```

5. Set Trigger Condition:
  - If took > 2000ms for 3 consecutive queries.

6. **Choose Action:**

- Send an alert via **Slack**:

```
{  
  "message": "Warning! Elasticsearch queries are taking more than 2s. Investigate query  
performance!"  
}
```

7. Save the rule.
-

### Example 3: Alert for Kubernetes Pod Network Latency

1. Go to Kibana → Stack Management → Rules & Connectors.
2. Click **Create Rule**.
3. Choose **Elasticsearch Query**.

4. Set:

- Name: **Kubernetes Network Latency Alert**
- Index: **k8s-metrics-\***.
- Query:

```
{  
  "query": {  
    "range": {  
      "latency": {  
        "gte": 500  
      }  
    }  
  }  
}
```

5. Set Trigger Condition:

- If latency > 500ms for more than **5 minutes**.

6. **Choose Action:**

- Send an **email alert** to DevOps team.
- Restart affected pods using **Kubernetes API calls**.

---

## 5. Real-time Monitoring Using Kibana

### Example: Monitoring Real-time Logs

1. Go to Kibana → Observability → Logs.
2. Select **Live Stream**.
3. Apply a filter:

response\_time > 1000

4. Observe real-time log entries.

---

## Step-by-Step Kibana Hands-on Lab

This hands-on lab will guide you through setting up **Kibana**, ingesting logs and metrics, visualizing performance data, and creating alerts.

---

### 1. Setup the ELK Stack (Elasticsearch, Logstash, Kibana)

#### 1.1 Install Elasticsearch

##### Step 1: Download and Install Elasticsearch

Run the following commands to download and extract Elasticsearch:

```
wget https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-8.4.0-linux-x86_64.tar.gz
tar -xzf elasticsearch-8.4.0-linux-x86_64.tar.gz
cd elasticsearch-8.4.0
```

##### Step 2: Configure Elasticsearch

Edit the configuration file config/elasticsearch.yml:

```
network.host: 0.0.0.0
http.port: 9200
cluster.name: "elk-cluster"
node.name: "node-1"
```

##### Step 3: Start Elasticsearch

```
./bin/elasticsearch
```

Verify it is running:

```
curl -X GET "http://localhost:9200"
```

---

#### 1.2 Install Kibana

##### Step 1: Download and Install Kibana

```
wget https://artifacts.elastic.co/downloads/kibana/kibana-8.4.0-linux-x86_64.tar.gz
tar -xzf kibana-8.4.0-linux-x86_64.tar.gz
cd kibana-8.4.0
```

## Step 2: Configure Kibana

Edit the configuration file config/kibana.yml:

```
server.port: 5601

server.host: "0.0.0.0"

elasticsearch.hosts: ["http://localhost:9200"]
```

## Step 3: Start Kibana

```
./bin/kibana
```

Access Kibana in your browser at **<http://localhost:5601/>**.

---

## 1.3 Install Logstash (Optional, for Log Parsing)

If you need to process logs before storing them in Elasticsearch:

```
wget https://artifacts.elastic.co/downloads/logstash/logstash-8.4.0-linux-x86_64.tar.gz
tar -xzf logstash-8.4.0-linux-x86_64.tar.gz
cd logstash-8.4.0
```

Configure **Logstash Pipeline** (logstash.conf):

```
input {
  file {
    path => "/var/log/app.log"
    start_position => "beginning"
  }
}

filter {
  grok {
    match => { "message" => "%{TIMESTAMP_ISO8601:timestamp} %{LOGLEVEL:loglevel}
%{GREEDYDATA:message}" }
  }
}

output {
  elasticsearch {
    hosts => ["http://localhost:9200"]
  }
}
```

```
    index => "app-logs"
  }
}
```

Run Logstash:

```
./bin/logstash -f logstash.conf
```

---

## 2. Ingest Logs and Metrics

### 2.1 Ship Application Logs Using Filebeat

#### Step 1: Install Filebeat

```
wget https://artifacts.elastic.co/downloads/beats/filebeat/filebeat-8.4.0-linux-x86_64.tar.gz
```

```
tar -xzf filebeat-8.4.0-linux-x86_64.tar.gz
```

```
cd filebeat-8.4.0
```

#### Step 2: Configure Filebeat

Edit filebeat.yml:

```
filebeat.inputs:
  - type: log
    paths:
      - "/var/log/app.log"
```

```
output.elasticsearch:
```

```
hosts: ["http://localhost:9200"]
```

#### Step 3: Start Filebeat

```
./filebeat -e
```

---

### 2.2 Collect System Metrics Using Metricbeat

#### Step 1: Install Metricbeat

```
wget https://artifacts.elastic.co/downloads/beats/metricbeat/metricbeat-8.4.0-linux-x86_64.tar.gz
```

```
tar -xzf metricbeat-8.4.0-linux-x86_64.tar.gz
```

```
cd metricbeat-8.4.0
```

---

## Step 2: Configure Metricbeat

Edit metricbeat.yml:

metricbeat.modules:

- module: system

metricsets: ["cpu", "memory", "load"]

period: 10s

output.elasticsearch:

hosts: ["http://localhost:9200"]

## Step 3: Start Metricbeat

./metricbeat -e

---

## 3. Create a Data View in Kibana

1. Open **Kibana** → **Stack Management** → **Data Views**.
  2. Click "**Create Data View**".
  3. Enter **app-logs-\*** as the pattern.
  4. Select the timestamp field (@timestamp).
  5. Click **Save**.
- 

## 4. Create Visualizations

### 4.1 Visualizing API Response Time

1. Go to **Kibana** → **Visualize**.
  2. Click **Create Visualization**.
  3. Choose **Line Chart**.
  4. Select **app-logs-\*** as the data source.
  5. Configure:
    - **Y-axis:** avg(response\_time)
    - **X-axis:** @timestamp
  6. Save and add to the **Dashboard**.
-

## 4.2 Detecting Frequent HTTP 500 Errors

1. Go to **Kibana** → **Visualize**.
  2. Click **Create Visualization**.
  3. Choose **Bar Chart**.
  4. Select **app-logs-\*** as the data source.
  5. Configure:
    - **Y-axis:** count(status\_code: 500)
    - **X-axis:** @timestamp
  6. Save the visualization.
- 

## 5. Create a Dashboard

1. Go to **Kibana** → **Dashboard**.
  2. Click **Create Dashboard**.
  3. Click **Add Existing Visualizations**:
    - API Response Time (Line Chart).
    - HTTP 500 Errors (Bar Chart).
  4. Save as **"Performance Monitoring Dashboard"**.
- 

## 6. Set Up Alerts

### 6.1 Alert for High API Response Times

1. Go to **Kibana** → **Stack Management** → **Rules & Connectors**.
2. Click **Create Rule**.
3. Choose **Elasticsearch Query**.
4. Set:
  - Name: **High API Latency Alert**
  - Index: **app-logs-\***
  - Query:

```
{  
  "query": {
```

```
"range": {
  "response_time": {
    "gte": 1000
  }
}
```

5. Set Trigger Condition:

- If response\_time > 1000ms for 5 minutes.

6. Choose Action:

- Send an **email or Slack notification**.

---

## 7. Real-time Log Monitoring

1. Go to Kibana → Observability → Logs.

2. Select **Live Stream**.

3. Apply filter:

response\_time > 1000

---

## 8. Advanced Monitoring

### 8.1 Elasticsearch Query Performance Monitoring

Enable **slow query logs** in Elasticsearch:

index.search.slowlog.threshold.query.warn: "2s"

Query slow logs in Kibana:

```
{
  "query": {
    "match": {
      "message": "took more than 2s"
    }
  }
}
```



---

## Correlating GC Logs, Heap Dumps, and JVM Metrics in Kibana

This step-by-step guide will help you correlate **GC logs, heap dumps, and JVM performance metrics** in Kibana to detect and analyze Java application performance issues such as **memory leaks, high GC pauses, and excessive heap consumption**.

---

### 1. Why Correlate GC Logs, Heap Dumps, and JVM Metrics?

Java applications running in production often suffer from **performance bottlenecks** related to:

- **Frequent Full GC pauses** causing application slowdowns.
- **High heap usage** leading to OutOfMemoryError (OOM).
- **Memory leaks** where objects are not being garbage collected.
- **Thread contention** affecting response times.

By **correlating GC logs, heap dumps, and JVM metrics in Kibana**, we can: ✓ Identify **slow GC cycles** and their impact on response times.

✓ Detect **heap memory spikes** and investigate retained objects.

✓ Monitor **CPU and thread activity** to find inefficiencies.

---

### 2. Setting Up Log Collection for JVM Monitoring

#### 2.1 Enabling GC Logs in Java

Java provides options to log **Garbage Collection (GC) activities**.

Modify the Java startup command:

```
java -Xlog:gc*:file=/var/log/gc.log:time,level,tags -jar myapp.jar
```

For Java 8:

```
java -XX:+PrintGCDetails -XX:+PrintGCDateStamps -Xloggc:/var/log/gc.log -jar myapp.jar
```

---

#### 2.2 Collecting GC Logs Using Filebeat

##### 1. Install Filebeat:

```
wget https://artifacts.elastic.co/downloads/beats/filebeat/filebeat-8.4.0-linux-x86_64.tar.gz
```

```
tar -xzf filebeat-8.4.0-linux-x86_64.tar.gz
```

```
cd filebeat-8.4.0
```

---

## 2. Configure Filebeat to Collect GC Logs

Edit filebeat.yml:

```
filebeat.inputs:

- type: log

paths:

- "/var/log/gc.log"

multiline.pattern: "^[0-9]{4}-[0-9]{2}-[0-9]{2}T[0-9]{2}:[0-9]{2}:[0-9]{2}"

multiline.negate: true

multiline.match: after


output.elasticsearch:

  hosts: ["http://localhost:9200"]

  index: "gc-logs"
```

## 3. Start Filebeat:

```
./filebeat -e
```

---

## 2.3 Collecting JVM Metrics Using Metricbeat

Metricbeat can collect **JVM heap usage, thread counts, and CPU load**.

### 1. Install Metricbeat:

```
wget https://artifacts.elastic.co/downloads/beats/metricbeat/metricbeat-8.4.0-linux-x86_64.tar.gz

tar -xzf metricbeat-8.4.0-linux-x86_64.tar.gz

cd metricbeat-8.4.0
```

### 2. Enable the Jolokia Module for JVM Monitoring:

```
metricbeat modules enable jolokia
```

### 3. Configure Metricbeat (metricbeat.yml):

```
metricbeat.modules:

- module: jolokia

  metricsets: ["jvm"]

  hosts: ["http://localhost:8778/jolokia"]
```

```
namespace: "jvm"
period: 10s
output.elasticsearch:
  hosts: ["http://localhost:9200"]
```

#### 4. Start Metricbeat:

```
./metricbeat -e
```

---

## 2.4 Capturing Heap Dumps for Memory Leak Analysis

### 1. Capture a Heap Dump When an OOM Occurs Modify Java options:

```
-XX:+HeapDumpOnOutOfMemoryError -XX:HeapDumpPath=/var/dumps/heapdump.hprof
```

### 2. Manually Capture a Heap Dump

```
jmap -dump:format=b,file=/var/dumps/heapdump.hprof <PID>
```

### 3. Ship Heap Dump Metadata to Elasticsearch Use Filebeat to collect heap dump metadata (heap-analyzer.log):

```
filebeat.inputs:
```

```
- type: log
```

```
paths:
```

```
- "/var/log/heap-analyzer.log"
```

---

## 3. Creating a Data View in Kibana

1. Go to Kibana → Stack Management → Data Views.
2. Click "Create Data View".
3. Enter gc-logs-\* as the pattern.
4. Select the timestamp field (@timestamp).
5. Click **Save**.

Repeat this for:

- jvm-metrics-\*
  - heap-metadata-\*
-

## 4. Creating Visualizations in Kibana

### 4.1 GC Pause Duration Over Time

1. **Go to Kibana → Visualize.**
  2. **Create a Line Chart.**
  3. Set:
    - **Y-axis:** avg(gc\_pause\_time).
    - **X-axis:** @timestamp.
  4. **Filter:** gc\_type: "Full GC".
  5. **Save the visualization.**
- 

### 4.2 Heap Usage Trend Analysis

1. **Go to Kibana → Visualize.**
  2. **Create an Area Chart.**
  3. Set:
    - **Y-axis:** avg(heap\_used).
    - **X-axis:** @timestamp.
  4. **Save and add to a Dashboard.**
- 

### 4.3 Correlating GC Logs and Heap Usage

1. **Create a New Dashboard.**
  2. Add:
    - **GC Pause Duration** (Line Chart).
    - **Heap Usage Trend** (Area Chart).
    - **Thread Count** (Metric Visualization).
  3. Apply filters:  
`gc_pause_time > 2s AND heap_used > 80%`
  4. Save the **"JVM Performance Dashboard"**.
- 

## 5. Setting Up Alerts

---

## 5.1 Alert for Long GC Pauses

1. Go to Kibana → Stack Management → Rules & Connectors.

2. Create a New Rule.

3. Select **Elasticsearch Query**.

4. Set:

- Name: **"High GC Pause Detected"**.

- Index: **gc-logs-\***.

- Query:

```
{
  "query": {
    "range": {
      "gc_pause_time": {
        "gte": 2000
      }
    }
  }
}
```

5. **Trigger Conditions:**

- If **GC Pause > 2s** for **3 occurrences**.

6. **Set Action:**

- Send **Slack alert:**

```
{
  "message": "High GC Pause detected! Investigate heap usage and GC logs."
}
```

---

## 5.2 Alert for High Heap Usage

1. Go to Kibana → Stack Management → Rules & Connectors.

2. Create a New Rule.

3. Select **Elasticsearch Query**.

4. Set:

- Name: **"Heap Usage Alert"**.
- Index: **jvm-metrics-\***.
- Query:

```
{  
  "query": {  
    "range": {  
      "heap_used": {  
        "gte": 85  
      }  
    }  
  }  
}
```

5. Trigger Conditions:

- If **heap usage > 85%** for **5 minutes**.

6. Set Action:

- Send an **email alert to DevOps team**.
-