Understanding & Diagnosing API Latency Spikes with Splunk

© Focus: Full-stack observability, log correlation, SPL deep dive, service tracing, RCA, and real examples

▶ 1. Architecture Context: Why Latency Spikes Are Hard to Trace

Modern microservices = asynchronous + distributed + multi-hop. One user request can span 5–15 services, sometimes across:

- AWS Lambda, EC2, K8s pods
- Managed DBs (Aurora, MongoDB Atlas)
- Queues (Kafka, SQS)
- External APIs (Stripe, Twilio)

Problem: Latency in one hop (say, Stripe call or thread lock in one container) causes end-to-end spikes.

Splunk becomes your microscope when every millisecond matters.

2. Log Design Best Practices for Splunk-Based Tracing

To trace latency:

O Use Consistent Structure for Logs

Log Format (JSON, consistent, indexed fields):

```
"timestamp": "2025-04-18T11:45:02Z",

"level": "INFO",

"service": "order-service",

"span_id": "a1b2c3",

"parent span id": "root",
```

```
"trace id": "TXN-001234",
 "duration ms": 1450,
 "operation": "placeOrder",
 "status": "success",
 "host": "order-node-1",
 "message": "Order placed successfully"
}
```

✓ Include:

- trace id: Unique per request
- span id: Per service operation
- duration ms: Local processing time
- host, container_id, or pod_id for infra mapping
- Optional: http status, db time, external api time

3. Step-by-Step Workflow for Diagnosing Latency

Step 1: Detect latency spikes using percentile SPLs

index=prod logs earliest=-15m

| stats perc50(duration_ms), perc90(duration_ms), perc99(duration_ms) by service

Check which service has **P90/P99 spike**, e.g., payment-service has perc90 > 1400ms.

Step 2: Identify spike pattern over time

index=prod logs service=payment-service

| timechart span=1m perc90(duration_ms)

- You now know:
 - When latency started
 - Whether it's constant, bursty, or increasing

Step 3: Drill down to affected transactions

index=prod_logs service=payment-service

| where duration_ms > 1000

| stats count, avg(duration_ms), max(duration_ms) by trace_id

| sort -max(duration_ms)



Output shows **trace IDs** where high latency happened.

Step 4: Trace full call chain of a single trace_id

Let's say trace id=TXN-001234

index=prod_logs trace_id="TXN-001234"

| sort _time

| table _time, service, operation, duration_ms, host, message



Example Output:

_time	service	operation	duration_ms	host	message
11:45:01.101	gateway	route	30	gw-1	Routed to checkout
11:45:01.134	auth-service	validateToken	50	auth- 1	Auth success
11:45:01.200	cart-service	getCart	90	cart-2	Cart fetched
11:45:01.300	inventory	checkItems	100	inv-2	Inventory confirmed
11:45:01.500	payment	chargeCard	1300	pay-1	Stripe API delay: 1200ms 🚹

11:45:02.900 noti	fy- sendSM	1S 60	sms-1	SMS sent
serv	ice			

Spike is in payment during Stripe API call.

Step 5: Check Deployment Events Around Spike

index=deployments service=payment-service

| table _time, version, deployed_by, commit_id

Correlate deploy time to spike time.

Step 6: Overlay Infra Metrics

Use metrics logs ingested into Splunk (from Telegraf, CloudWatch, Prometheus exporters).

Overlay CPU & Memory of affected host

index=infra_metrics host="pay-1"

| timechart span=1m avg(cpu_utilization), avg(memory_used)

High CPU or memory pressure may cause GC pauses, IO delays.

Step 7: Heatmap of Latency Distribution

index=prod_logs service=payment-service

| bucket duration_ms span=100

| stats count by duration ms

sort duration_ms

Helps build latency histogram.

Step 8: Correlate with Error, GC, Thread Logs

6 Find Full GC Events

index=gc_logs host="pay-1" message="Full GC"

| table _time, message, duration

Correlate Full GC (1.2s) with your trace latency!

Detect Thread Blocking

index=thread_logs host="pay-1" message="blocked"

| stats count by thread_id

Step 9: Automate RCA with transaction or stats

index=prod_logs service=payment-service

| transaction trace id startswith="Payment request started" endswith="Payment complete"

| eval total time = duration * 1000

| where total_time > 1000

| table trace id, total time, host

Or alternatively:

index=prod_logs service=payment-service

stats sum(duration ms) as total duration by trace id

| where total duration > 1000

Step 10: Look for Common API/DB Bottlenecks

© External API (Stripe, Razorpay, etc.)

index=prod logs service=payment-service message="Calling Stripe"

| stats avg(duration ms), count by api url

DB Query Analysis

index=prod_logs service=payment-service message="Executing query"

| stats avg(duration_ms), count by sql_query

| sort -avg(duration_ms)

RCA Matrix Cheat Sheet

Suspected Cause	SPL Example / Log Signal		
External API delay	message="Stripe API" + duration > 1000		
Thread Pool Saturation	message="ThreadPoolExhausted"		
GC Pause	index=gc_logs Full GC		
N+1 Query Pattern	High # of similar SQLs per trace		
Load Balancer Delay	Logs from gateway show high queue_time_ms		
Network I/O delay	message="HTTP timeout"		
CPU pressure	cpu_utilization > 85%		
Memory leak / GC churn	Rising GC duration + heap_used across traces		

Final Dashboards You Must Have

Dashboard Panel SPL Summary P90 Latency Over Time timechart span=1m perc90(duration ms) `stats avg(duration_ms) by service Top 5 Slow Services Slowest Transactions (trace_id) where duration_ms > 1000 + stats by trace_id **GC Events Timeline** `index=gc_logs Full GC message="Stripe API" + stats by api_url External API Call Latency Heatmap of duration buckets `bucket duration_ms span=100 Infra overlays (CPU/Memory) timechart avg(cpu_utilization), avg(memory_used)

Advanced Use Case: Automated Anomaly Detection

Use **Splunk Machine Learning Toolkit** or SPL thresholds:

index=prod logs service=payment-service

| timechart span=1m avg(duration_ms) as duration

| anomalydetection duration

→ Setup alerts when P90 > 1000ms or Stripe API > 1200ms using savedsearch.

Summary – What to Do When You See Latency Spikes

- 1. Check percentiles (p90, p99) across services
- 2. **Drill down by trace ID** using duration_ms > threshold
- 3. Trace full call path → locate spike
- 4. Correlate with deployments, infra logs, GC, thread logs
- 5. Use SPL to find common patterns in external APIs or DB
- 6. Automate dashboards and alerts for future detection